

Building BT Network

by

Zhihao DAI

Yunsong ZHANG

Huijing LEI

Changrong CHEN

Yan HUANG

**19COP502 Building Secure Networks
Lab Report**

Loughborough University

© BT NETWORK 2019

Nov. 2019

Abstract

This report describes the building process of BT Network, a small version of Tier-2 Internet Service Provider. To first define the architecture of the network, a network diagram is drawn, in which physical connections are established and IP addresses are assigned. Then, connectivity within the network and with outside networks are made possible through internal routing protocol Intermediate System to Intermediate System (IS-IS) and external Border Gateway Protocol (BGP). In addition, services such as World Wide Web (WWW), Domain Name System (DNS) and Email are provided in the network. At last, main conclusions and further improvements are discussed.

Contents

Abstract	i
List of Figures	iv
List of Tables	vi
List of Listings	vii
1 Introduction	1
1.1 Network Services	1
1.2 Business Relationships with Neighbouring ISPs	2
1.2.1 Provider: Central Network	2
1.2.2 Peer: Virgin Network	3
1.2.3 Customer: DT Network	3
1.3 Roles of Network Components	3
1.3.1 Routers	3
1.3.2 Laptops	3
1.4 Organisation of the Report	4
2 Network Architecture	6
2.1 Network Diagram	6
2.1.1 Description	6
2.1.2 IP Addresses and Interfaces	8
2.2 IP Addresses of Interfaces	11
2.2.1 Implementation	11
2.2.2 Evaluation	13
3 Routing Protocols in the Network	16
3.1 Intra-domain Routing Protocol: IS-IS	16

CONTENTS

3.1.1	Design	16
3.1.2	Loopback Addresses and NSAP for Routers	16
3.1.3	Implementation	18
3.1.4	Evaluation	19
3.1.5	Commentary	25
3.2	Inter-domain Routing Protocol: BGP	27
3.2.1	Design	27
3.2.2	Routing Policies	28
3.2.3	Implementation	28
3.2.4	Evaluation	30
3.2.5	Commentary	39
4	Applications in the Network	41
4.1	Secure Remote Access to Routers through SSH	41
4.1.1	Design	41
4.1.2	Implementation	41
4.1.3	Evaluation	42
4.1.4	Commentary	43
4.2	Domain Name System Service	44
4.2.1	Design	44
4.2.2	Implementation	44
4.2.3	Evaluation	47
4.2.4	Commentary	51
4.3	World Wide Web Service	53
4.3.1	Design	53
4.3.2	Implementation	53
4.3.3	Evaluation	54
4.4	Email Service	55
4.4.1	Design	55
4.4.2	Implementation	55
4.4.3	Evaluation	55
4.4.4	Commentary	55
5	Discussion	58
5.1	Conclusions	58
5.2	Further Work	58

CONTENTS

6 Contributions	59
6.1 Group Leader: Zhihao DAI	59
6.2 Technical Director: Yunsong ZHANG	59
6.3 Network Engineer: Huijing LEI	59
6.4 Network Engineer: Changrong CHEN	59
6.5 Network Engineer: Yan HUANG	60
References	61
A Login Details	62
A.1 Laptops	62
A.2 Routers	62
B Routers Configuration	63
B.1 Router 1 Configuration	63
B.2 Router 2 Configuration	67
B.3 Router 3 Configuration	71
C Laptops Configuration	75
C.1 Laptop 1 Configuration	75
C.2 Laptop 2 Configuration	76
C.3 Laptop 3 Configuration	77
D DNS Configuration	78
D.1 Primary DNS Server (Laptop 3)	78
D.2 Secondary DNS Server (Laptop 1)	80
E Email Configuration	82

List of Figures

1.1	Business Relationships of BT Network with Neighbouring ISPs.	2
1.2	Roles of Main Components in BT Network.	4
2.1	Full Network Diagram of BT Network.	7
2.2	Physical Connections within BT Network.	8
2.3	Successful Assignment of IPv4 Addresses to Routers' Interfaces.	14
2.4	Successful Assignment of IPv6 Addresses to Routers' Interfaces.	14
2.5	Successful Assignment of IPv4 and IPv6 Addresses to Laptops' Interfaces. . .	15
3.1	Design of IS-IS Protocol in BT Network.	17
3.2	Tracing IPv4 Routes between Laptops using <code>tracert</code>	20
3.3	Tracing IPv6 Routes between Laptops using <code>tracert</code>	21
3.4	Tracing IPv4 Routes from Laptop 2 (BT002) to Laptop 3 (BT003) when Physical Connection between Router 1 (BT-R001) and Router 2 (BT-R002) is broken.	22
3.5	Tracing IPv6 Routes from Laptop 2 (BT002) to Laptop 3 (BT003) when Physical Connection between Router 1 (BT-R001) and Router 2 (BT-R002) is broken.	22
3.6	IPv4 Routes to Other Subnets on All 3 Routers Respectively using <code>show ip route</code>	23
3.7	IPv6 Routes to Other Subnets on All 3 Routers Respectively using <code>show ipv6 route</code>	24
3.8	IP Routes to Other Subnets on Router 1 when Physical Connection between Router 1 (BT-R001) and Router 2 (BT-R002) is broken.	26
3.9	Design of BGP Protocol in BT Network.	27
3.10	IPv4 Routes Collected through BGP Protocols on All 3 Routers using <code>show bgp</code>	31

LIST OF FIGURES

3.11 IPv6 Routes Collected through BGP Protocols on Routers using <code>show bgp</code> <code>ipv6</code>	32
3.12 Tracing IPv4 Routes to Central Network on Laptop 1 (BT001) using <code>traceroute</code> .	33
3.13 Tracing IPv4 Routes to DT Network on Laptop 1 (BT001) using <code>traceroute</code> .	35
3.14 Tracing IPv4 Routes to Virgin Network on Laptop 1 (BT001) using <code>traceroute</code> .	37
3.15 Tracing IPv4 Routes to Other Networks on Laptop 1 (BT001) using <code>traceroute</code> .	38
3.16 Tracing IPv4 Routes to Virgin Network on Laptop 1 (BT001) using <code>traceroute</code> When Direct Physical Connection Is Broken	40
4.1 Successful remote SSH access to all 3 routers from Laptop 1 (BT001).	42
4.2 Querying DNS Records for <code>bt.lboro</code> and <code>www.bt.lboro</code> on Laptop 2.	48
4.3 Querying DNS Records for sub-domains of <code>bt.lboro</code> on Laptop 2.	48
4.4 Querying DNS AAAA Records for sub-domains of <code>bt.lboro</code> on Laptop 2.	49
4.5 Querying Reverse DNS Records for All 3 Laptops on Laptop 2.	49
4.6 Connectivity to <code>bt.lboro</code> using <code>ping</code> on all 3 Laptops.	50
4.7 Querying DNS Records for <code>bt.lboro</code> on Laptop 2 When Primary DNS Server (Laptop 3) Is Down.	51
4.8 Contents of HTML File Named <code>index.html</code>	53
4.9 Web Service Provided at <code>http://bt.lboro</code>	54
4.10 Important Configuration Steps for <code>exim4</code>	56
4.11 Email Can Be Seen Arrived on Both Sides.	57

List of Tables

2.1	Allocation of IPv4 and IPv6 Addresses to Subnets in BT Network.	10
2.2	Interfaces for Each Physical Connection and Corresponding IPv4 and IPv6 Addresses.	10
3.1	IP Loopback Addresses and NSAP for Routers in BT Network.	17

List of Listings

4.1	Contents of Forward DNS Configuration File Located at <code>/etc/bind/db.bt.lboro</code> on Primary DNS Server (Laptop 3).	45
4.2	Contents of Reverse DNS Configuration File Located at <code>/etc/bind/db.23</code> on Primary DNS Server (Laptop 3).	46
B.1	Contents of Configuration on Router 1 (BT-R001).	63
B.2	Contents of Configuration on Router 2 (BT-R002).	67
B.3	Contents of Configuration on Router 3 (BT-R003).	71
C.1	Contents of Interfaces Network Configuration File Located at <code>/etc/network/interfaces</code> on Laptop 1 (BT001).	75
C.2	Contents of Interfaces Network Configuration File Located at <code>/etc/network/interfaces</code> on Laptop 2 (BT002).	76
C.3	Contents of Interfaces Network Configuration File Located at <code>/etc/network/interfaces</code> on Laptop 3 (BT003).	77
D.1	Contents of DNS Configuration File Located at <code>/etc/bind/named.conf.options</code> on Primary DNS Server (Laptop 3).	78
D.2	Contents of Local DNS Configuration File Located at <code>/etc/bind/named.conf.local</code> on Primary DNS Server (Laptop 3).	79
D.3	Contents of DNS Configuration File Located at <code>/etc/bind/named.conf.options</code> on Secondary DNS Server (Laptop 1).	80
D.4	Contents of Local DNS Configuration File Located at <code>/etc/bind/named.conf.local</code> on Secondary DNS Server (Laptop 1).	80
E.1	Contents of Generated Email Configuration File Located at <code>/var/lib/exim4/config.autogenerated</code> on Email Server (Laptop 3).	82

Chapter 1

Introduction

In this lab, our team sets out to build BT Network, a small version of Tier-2 Internet Service Provider (ISP) located at Haslegrave Building, from scratch. Despite its limitations in terms of size and Internet access, we can proudly attest that BT Network is one of the leading providers at Haslegrave Building.

BT network is an Autonomous System (AS) as a whole and the AS Number is 2030. Its domain name is `bt.lboro`.

1.1 Network Services

Our network provides the following services to each of our individual customers.

- IP addressing with a guaranteed range of 14 host addresses allocated from 23.0.0.0/8 (IPv4) and 2001:2300::/32 (IPv6) blocks.
- Intra-domain Internet connection with Intermediate System to Intermediate System (IS-IS) routing protocol.
- Inter-domain Internet connection with Border Gateway Protocol (BGP).
- A reliable Domain name System (DNS) service with duplicated servers under domain `bt.lboro`.
- A World Wide Web (WWW) service located at `http://bt.lboro/`.
- An Email service at `bt.lboro`.

For neighbouring ISPs who is a customer in our business relationships (see Section 1.2.3), we provide the following services.

- Internet connection to our domain as well as all the others’.

In addition, we provide secure remote access to our routers through Secure Shell (SSH) protocol on one of our laptops for administrative purposes.

1.2 Business Relationships with Neighbouring ISPs

BT Network has three immediate neighbouring ISPs and it’s important to form business relationships with all three of them in order to gain economic benefits. The external routing policies of BGP protocol for each outside network are determined by its business relationship with us (see Section 3.2 for details). Business relationships with neighbouring ISPs are shown in Figure 1.1 and elaborated in the following.

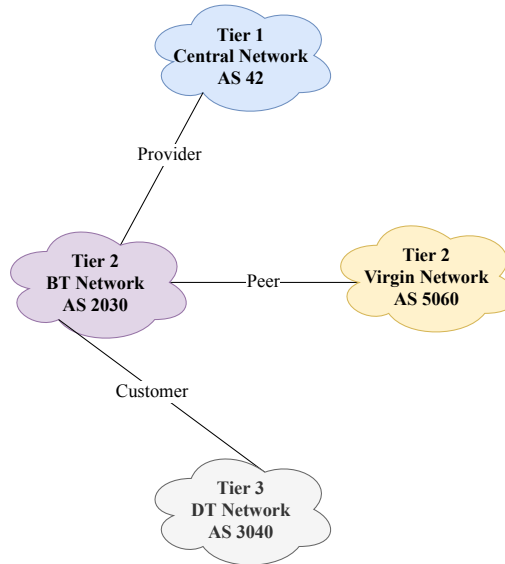


Figure 1.1: Business Relationships of BT Network with Neighbouring ISPs.

1.2.1 Provider: Central Network

Since BT Network is a Tier-2 ISP, it need to be connected to a Tier-1 ISP to gain broader Internet connection. Therefore, BT is connected to Central Network, a Tier-1 ISP, as a customer which makes it a network provider for BT.

1.2.2 Peer: Virgin Network

BT Network forms a Peer relationship with Virgin Network. This allows Virgin Network to connect to BT Network at zero cost and vice versa.

1.2.3 Customer: DT Network

BT Network forms a Provider-Customer relationship with DT Network, in which BT is the provider and DT is the customer. In other words, DT gains access to the broader Internet through BT at a cost.

1.3 Roles of Network Components

There are 6 physical components in our network in total, of which 3 are Cisco routers and the other 3 are TOSHIBA laptops. Each component plays an important role in the network as shown in Figure 1.2.

1.3.1 Routers

In terms of connection, each router is attached with one customer subnet and thus providing Internet service to one customer. Router 1 (BT-R001) is not physically connected to any outside network. Router 2 (BT-R002) is connected to DT Network and Router 3 (BT-R003) is connected to Virgin Network and Central Network through cables.

In terms of routing, all routers are Level-1 routers in intra-domain IS-IS routing protocol. In BGP routing protocol, Router 1 (BT-R001) acts as an Internal BGP (IBGP) router while Router 2 (BT-R002) and Router 3 (BT-R003) act as External BGP (EBGP) routers.

1.3.2 Laptops

All laptops are running a Ubuntu 16.04 system. Each of them is connected to a customer subnet through a cable. In terms of services, Laptop 1 (BT001) provides DNS service for `bt.lboro` as a secondary DNS server and WWW service at `http://bt.lboro`. It also acts as a secure SSH access point to routers for administrative purposes. Laptop 2 (BT002) doesn't provide any service and thus acts as a normal user in the network. Laptop 3 (BT003) provides DNS service for `bt.lboro` as a primary DNS server. In addition, it hosts Email service at `bt.lboro`.

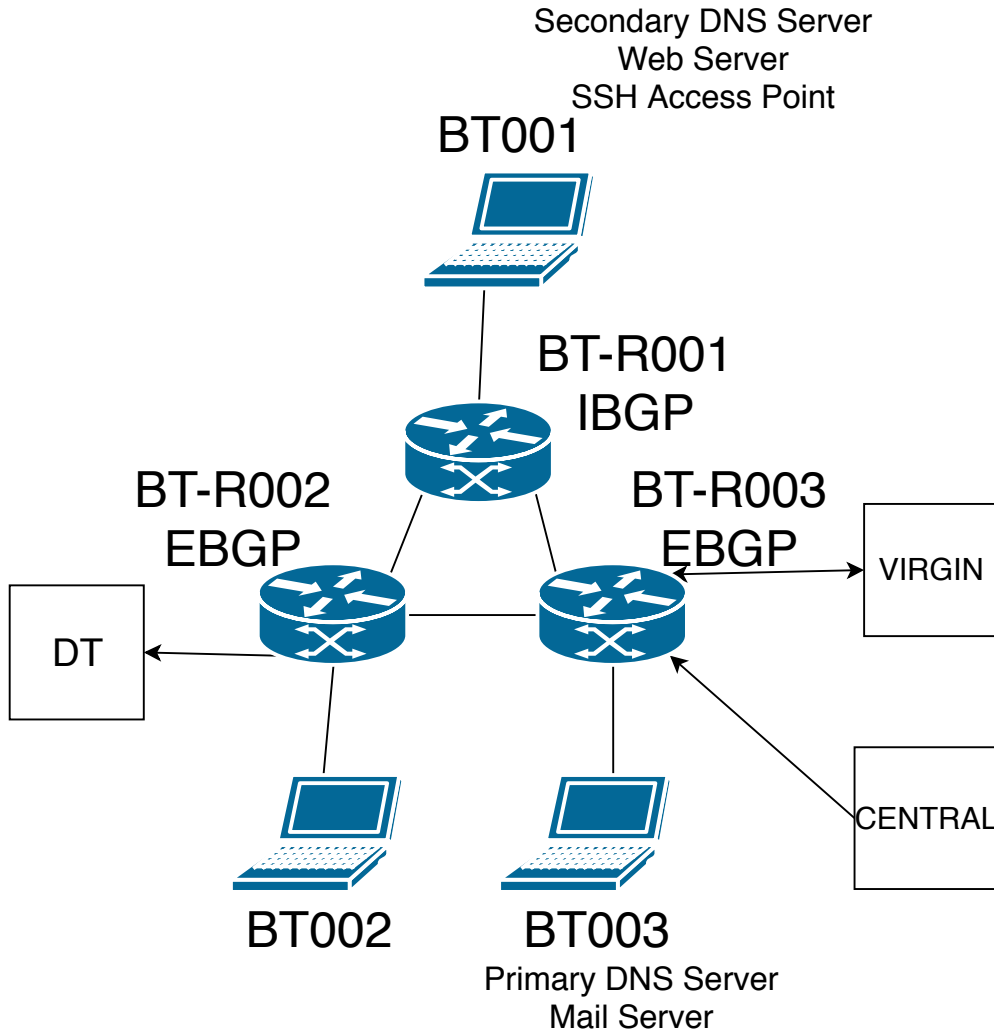


Figure 1.2: Roles of Main Components in BT Network.

1.4 Organisation of the Report

The report is organised as follows. The architecture of BT Network on the network layer is described in Chapter 2. A Network Diagram involving all physical components and connections is drawn in the chapter. Then, IP addresses for interfaces in the network are carefully allocated and configured.

In order to allow packets to be forwarded within and outside the network, proper intra-domain and inter-domain routing protocols are implemented and tested in Chapter 3.

In Chapter 4, we move up to the application layer and set up various services in the

network as listed in Section 1.1.

Main conclusions drawn from the building process and possible further improvements are discussed in Chapter 5. At last, a summary of contributions for each team member is presented in Chapter 6.

Although the report does not necessarily reflect the actual order of steps in our building process (eg. remote SSH access was set up before BGP), readers can be assured that all results presented can be reproduced by following the natural order of the report.

To ensure readability, rationale behind important decisions made, problems we encountered and their respective solutions, alternative ways of configurations (if any) as well as reflective commentary are documented along with each step of implementation in the report.

Chapter 2

Network Architecture

2.1 Network Diagram

2.1.1 Description

A full diagram of BT network is shown in Figure 2.1. There are 3 routers in the network, whose names are BT-R001, BT-R002 and BT-R003 respectively, connected to each other. Each connection forms a Router-Router subnet with only 2 interfaces.

On the other hand, each router is connected with a laptop separately named as BT001, BT002 and BT003 and thus forms a Router-Laptop subnet. A customer of BT Network is assigned with a Router-Laptop subnet and has a minimum of 10 host IP addresses.

To connect to neighbouring ISPs, a Router-Neighbour subnet is formed for each connection. Concretely, Router 2 (BT-R002) is connected to one of DT Network's routers while Router 3 (BT-R003) is connected to one of Virgin Network's and Central Network's routers separately.

All of the above connections are through physical cables.

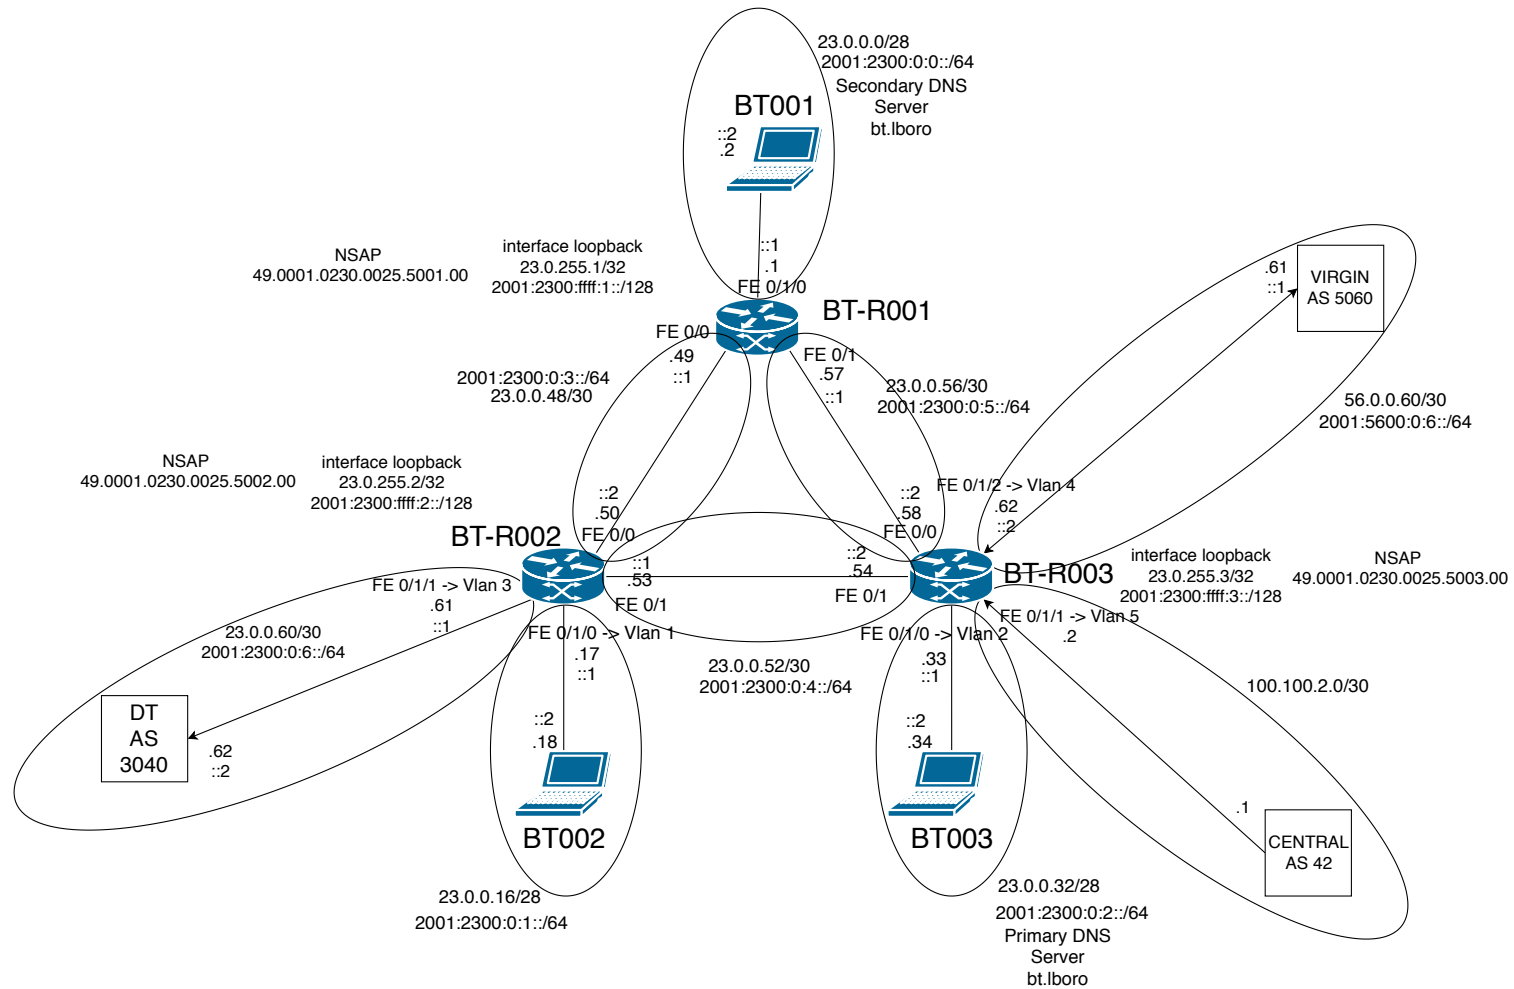


Figure 2.1: Full Network Diagram of BT Network.

2.1.2 IP Addresses and Interfaces

An IPv4 address range of 23.0.0.0/8 and IPv6 address range of 2001:2300::/32 are allocated to BT Network, which are further divided into sub-ranges to be allocated to each subnet.

For IPv4 addressing, a prefix of n is needed for a subnet that demands X host addresses, where n is an integer that satisfies $2^{32-n} - 2 \geq X$ and $n \leq 32$. For our lab, the maximum value for prefix is used in order to minimize the size of each subnet and reserve address space for future customers. **However, it's also possible to use a larger value for each Router-Laptop subnet in order to maximize the size of the subnet, given that the number of customers (in this case, 3) is fixed.**

In BT Network, the prefix for each Router-Router and Router-Neighbour subnet is 30 while the prefix for each Router-Laptop subnet is 28. In other words, each Router-Router and Router-Neighbour subnet has 2 guaranteed IPv4 host addresses while each Router-Laptop subnet has 14 guaranteed IPv4 host addresses. During address block allocation, larger subnet is being considered before smaller one reduce the number of block segments.

For IPv6 addressing, however, each subnet has a fixed prefix of 64 to ensure that each interface in the subnet has a unique address. The full details of IP address allocation is shown in Table 2.1.



Figure 2.2: Physical Connections within BT Network.

In terms of interfaces, there are 3 Ethernet interfaces (**FastEthernet0/0**, **FastEthernet0/1** and **FastEthernet0/1/0**) on Router 1, each of which can be assigned with an IP address. On Router 2 and 3, however, there are 6 Ethernet interfaces each and only 2 of them (**FastEthernet0/0** and **FastEthernet0/1**) can be directly assigned with IP addresses. The remaining 4 interfaces are link layer interfaces and thus does not possess any IP address. To be assigned with an IP address, such an interface need to be assigned to an Virtual LAN (VLAN) to which the address is actually assigned.

Router-Router connections are established through either **FastEthernet0/0** or **FastEthernet0/1** interfaces on both ends while Router-Laptop and Router-Neighbour are through one of the remaining interfaces on the router end. Since both interfaces are on the left-hand side of each router and such arrangement helps distinguishing between Router-Router connections and others easily as shown in Figure 2.2. Interfaces of both ends for each connection as well as their corresponding IP addresses are detailed in Table 2.2.

Subnet	IPv4 Address / Prefix	IPv4 Address Range	IPv6 Address / Prefix	IPv6 Address Range
BT-R001 - BT001	23.0.0.0/28	23.0.0.1 - 23.0.0.14	2001:2300:0:0::/64	2001:2300:0:0::1 - 2001:2300:0:0:ffff:ffff:ffff:fffe
BT-R002 - BT002	23.0.0.16/28	23.0.0.17 - 23.0.0.30	2001:2300:0:1::/64	2001:2300:0:1::1 - 2001:2300:0:1:ffff:ffff:ffff:fffe
BT-R003 - BT003	23.0.0.32/28	23.0.0.33 - 23.0.0.62	2001:2300:0:2::/64	2001:2300:0:2::1 - 2001:2300:0:2:ffff:ffff:ffff:fffe
BT-R001 - BT-R002	23.0.0.48/30	23.0.0.49 - 23.0.0.50	2001:2300:0:3::/64	2001:2300:0:3::1 - 2001:2300:0:3:ffff:ffff:ffff:fffe
BT-R002 - BT-R003	23.0.0.52/30	23.0.0.53 - 23.0.0.54	2001:2300:0:4::/64	2001:2300:0:4::1 - 2001:2300:0:4:ffff:ffff:ffff:fffe
BT-R001 - BT-R003	23.0.0.56/30	23.0.0.57 - 23.0.0.58	2001:2300:0:5::/64	2001:2300:0:5::1 - 2001:2300:0:5:ffff:ffff:ffff:fffe
BT-R002 - DT	23.0.0.60/30	23.0.0.61 - 23.0.0.62	2001:2300:0:6::/64	2001:2300:0:6::1 - 2001:2300:0:6:ffff:ffff:ffff:fffe
BT-R003 - Virgin	56.0.0.60/30	56.0.0.61 - 56.0.0.62	2001:5600:0:6::/64	2001:5600:0:6::1 - 2001:5600:0:6:ffff:ffff:ffff:fffe
BT-R003 - Central	100.100.2.0/30	100.100.2.1 - 100.100.2.2		

Table 2.1: Allocation of IPv4 and IPv6 Addresses to Subnets in BT Network.

Connection	Interface 1	IPv4 Address	IPv6 Address	Interface 2	IPv4 Address	IPv6 Address
BT-R001 - BT001	BT-R001: FastEthernet0/1/0	23.0.0.1	2001:2300:0:0::1	BT001: eth0	23.0.0.2	2001:2300:0:0::2
BT-R002 - BT002	BT-R002: FastEthernet0/1/0 -> Vlan 1	23.0.0.17	2001:2300:0:1::1	BT002: eth0	23.0.0.18	2001:2300:0:1::2
BT-R003 - BT003	BT-R003: FastEthernet0/1/0 -> Vlan 2	23.0.0.33	2001:2300:0:2::1	BT003: eth0	23.0.0.34	2001:2300:0:2::2
BT-R001 - BT-R002	BT-R001: FastEthernet0/0	23.0.0.49	2001:2300:0:3::1	BT-R002: FastEthernet0/0	23.0.0.50	2001:2300:0:3::2
BT-R002 - BT-R003	BT-R002: FastEthernet0/1	23.0.0.53	2001:2300:0:4::1	BT-R003: FastEthernet0/1	23.0.0.54	2001:2300:0:4::2
BT-R001 - BT-R003	BT-R001: FastEthernet0/1	23.0.0.57	2001:2300:0:5::1	BT-R003: FastEthernet0/0	23.0.0.58	2001:2300:0:5::2
BT-R002 - DT	BT-R002: FastEthernet0/1/1 -> Vlan 3	23.0.0.61	2001:2300:0:6::1	DT	23.0.0.62	2001:2300:0:6::2
BT-R003 - Virgin	BT-R003: FastEthernet0/1/2 -> Vlan 4	56.0.0.62	2001:5600:0:6::2	Virgin	56.0.0.61	2001:5600:0:6::1
BT-R003 - Central	BT-R003: FastEthernet0/1/1 -> Vlan 5	100.100.2.2		Central	100.100.2.1	

Table 2.2: Interfaces for Each Physical Connection and Corresponding IPv4 and IPv6 Addresses.

2.2 IP Addresses of Interfaces

Assigning IP addresses[1][2] to interfaces should be the first step in building BT Network since all network services listed in 1.1 cannot operate without IP addresses. The assignment of IP addresses in Table 2.1 and 2.2 is implemented.

2.2.1 Implementation

2.2.1.1 Routers

For Router 1 (BT-R001), IP addresses are assigned directly to physical interfaces as all interfaces are network layer interfaces.

```
1 int fa0/0
2 ip address 23.0.0.49 255.255.255.252
3 ipv6 address 2001:2300:0:3::1/64
4 no shutdown
5
6 int fa0/1
7 ip address 23.0.0.57 255.255.255.252
8 ipv6 address 2001:2300:0:5::1/64
9 no shutdown
10
11 int fa0/1/0
12 ip address 23.0.0.1 255.255.255.240
13 ipv6 address 2001:2300:0:0::1/64
14 no shutdown
```

For Router 2 (BT-R002), however, only 2 interfaces (FastEthernet0/0 and FastEthernet0/1) each are network layer interfaces. The remaining 4 interfaces are link layer interfaces and need to be assigned to an VLAN separately so where an IP address can be assigned.

```
1 int fa0/0
2 ip address 23.0.0.50 255.255.255.252
3 ipv6 address 2001:2300:0:3::2/64
4 no shutdown
5
6 int fa0/1
7 ip address 23.0.0.53 255.255.255.252
8 ipv6 address 2001:2300:0:4::1/64
9 no shutdown
10
11 vlan 1
12 int fa0/1/0
13 switchport mode access
14 switchport access vlan 1
```

```
15 int vlan 1
16 ip address 23.0.0.17 255.255.255.240
17 ipv6 address 2001:2300:0:1::1/64
18 no shutdown
19
20 vlan 3
21 int fa0/1/1
22 switchport mode access
23 switchport access vlan 3
24 int vlan 3
25 ip address 23.0.0.61 255.255.255.252
26 ipv6 address 2001:2300:0:6::1/64
27 no shutdown
```

Similarly, IP addresses are assigned to Router 3 (BT-R003).

```
1 int fa0/0
2 ip address 23.0.0.58 255.255.255.252
3 ipv6 address 2001:2300:0:5::2/64
4 no shutdown
5
6 int fa0/1
7 ip address 23.0.0.54 255.255.255.252
8 ipv6 address 2001:2300:4::2/64
9 no shutdown
10
11 vlan 2
12 int fa0/1/0
13 switchport mode access
14 switchport access vlan 2
15 int vlan 2
16 ip address 23.0.0.33 255.255.255.240
17 ipv6 address 2001:2300:0:2::1/64
18 no shutdown
19
20 vlan 4
21 int fa0/1/2
22 switchport mode access
23 switchport access vlan 4
24 int vlan 4
25 ip address 56.0.0.62 255.255.255.252
26 ipv6 address 2001:5600:0:6::2/64
27 no shutdown
28
29 vlan 5
30 int fa0/1/1
31 switchport mode access
32 switchport access vlan 5
```

```
33 int vlan 5
34 ip address 100.100.2.2 255.255.255.252
35 no shutdown
```

2.2.1.2 Laptops

Unlike routers, the IP assignment of laptops' interfaces are more aligned since all laptops are of the same model. On Laptop 1 (BT001), the file `/etc/network/interface` is edited as follows before rebooting to apply changes. Similarly, IP address can be assigned to Laptop 2's (BT002) and Laptop 3's (BT003) interfaces respectively. Full configuration for all laptops are detailed in Appendix C.

```
1 auto lo
2 iface lo inet loopback
3 auto eth0
4 iface eth0 inet static
5 address 23.0.0.2
6 netmask 255.255.255.240
7 gateway 23.0.0.1
8
9 iface eth0 inet6 static
10 address 2001:2300::2
11 netmask 64
12 gateway 2001:2300::1
```

2.2.2 Evaluation


On all 3 routers, the implementation of IP assignments is evaluated using `show ip int brief` and `show ipv6 int brief`. These 2 commands show IPv4 and IPv6 addresses of interfaces on routers respectively. Successful assignment of IP addresses to routers' interfaces is evident from Figure 2.3 and 2.4.

On the laptops' side, the implementation of IP assignments is evaluated using `ifconfig` command. Successful assignment of IP addresses to laptops' interfaces is evident from Figure 2.5.


CHAPTER 2. NETWORK ARCHITECTURE



(a) Router 1 (BT-R001)




(b) Router 2 (BT-R002)




(c) Router 3 (BT-R003)

Figure 2.3: Successful Assignment of IPv4 Addresses to Routers' Interfaces.



(a) Router 1 (BT-R001)



(b) Router 2 (BT-R002)



(c) Router 3 (BT-R003)

Figure 2.4: Successful Assignment of IPv6 Addresses to Routers' Interfaces.

```

4 23.0.0.34 (23.0.0.34) 1.432 ms 1.384 ms 1.362 ms
root@BT001:/home/bt001# ifconfig
eth0
    Link encap:Ethernet  HWaddr 00:1e:68:d0:d5:1a
    inet addr:23.0.0.2  Bcast:23.0.0.15  Mask:255.255.255.240
    inet6 addr: fe80::21e:68ff:fed0:d51a/64 Scope:Link
    UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
    RX packets:407 errors:0 dropped:0 overruns:0 frame:0
    TX packets:646 errors:0 dropped:0 overruns:0 carrier:0
    collisions:0 txqueuelen:1000
    RX bytes:41002 (41.0 KB)  TX bytes:51773 (51.7 KB)
    Interrupt:16

lo
    Link encap:Local Loopback
    inet addr:127.0.0.1  Mask:255.0.0.0
    inet6 addr: ::1/128 Scope:Host
    UP LOOPBACK RUNNING  MTU:65536  Metric:1
    RX packets:139 errors:0 dropped:0 overruns:0 frame:0
    TX packets:139 errors:0 dropped:0 overruns:0 carrier:0
    collisions:0 txqueuelen:0
    RX bytes:9727 (9.7 KB)  TX bytes:9727 (9.7 KB)

root@BT001:/home/bt001#

```

(a) Laptop 1 (BT001)

```

root@BT002:/# ifconfig
eth0
    Link encap:Ethernet  HWaddr 00:1e:68:d0:d5:1a
    inet addr:23.0.0.18  Bcast:23.0.0.31  Mask:255.255.255.240
    inet6 addr: fe80::21e:68ff:fed0:d51a/64 Scope:Link
    UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
    RX packets:103 errors:0 dropped:0 overruns:0 frame:0
    TX packets:103 errors:0 dropped:0 overruns:0 carrier:0
    collisions:0 txqueuelen:1000
    RX bytes:3636 (3.6 KB)  TX bytes:7240 (7.2 KB)
    Interrupt:16

lo
    Link encap:Local Loopback
    inet addr:127.0.0.1  Mask:255.0.0.0
    inet6 addr: ::1/128 Scope:Host
    UP LOOPBACK RUNNING  MTU:65536  Metric:1
    RX packets:125 errors:0 dropped:0 overruns:0 frame:0
    TX packets:125 errors:0 dropped:0 overruns:0 carrier:0
    collisions:0 txqueuelen:0
    RX bytes:19970 (19.7 KB)  TX bytes:19970 (19.7 KB)

root@BT002:/#

```

(b) Laptop 2 (BT002)

```

root@BT003:/home/bt003# ifconfig
eth0
    Link encap:Ethernet  HWaddr 00:1e:68:d0:d5:1a
    inet addr:23.0.0.34  Bcast:23.0.0.47  Mask:255.255.255.240
    inet6 addr: fe80::21e:68ff:fed0:d51a/64 Scope:Link
    UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
    RX packets:181 errors:0 dropped:0 overruns:0 frame:0
    TX packets:181 errors:0 dropped:0 overruns:0 carrier:0
    collisions:0 txqueuelen:1000
    RX bytes:15738 (15.7 KB)  TX bytes:11698 (11.6 KB)
    Interrupt:16

lo
    Link encap:Local Loopback
    inet addr:127.0.0.1  Mask:255.0.0.0
    inet6 addr: ::1/128 Scope:Host
    UP LOOPBACK RUNNING  MTU:65536  Metric:1
    RX packets:127 errors:0 dropped:0 overruns:0 frame:0
    TX packets:127 errors:0 dropped:0 overruns:0 carrier:0
    collisions:0 txqueuelen:0
    RX bytes:20202 (20.2 KB)  TX bytes:20202 (20.2 KB)

root@BT003:/home/bt003#

```

(c) Laptop 3 (BT003)

Figure 2.5: Successful Assignment of IPv4 and IPv6 Addresses to Laptops' Interfaces.

Chapter 3

Routing Protocols in the Network

3.1 Intra-domain Routing Protocol: IS-IS

3.1.1 Design

Due to the limitation of the static routing, routers cannot find alternative paths if a set path is broken and thus a new path need to be set manually. In contrast, dynamic routing always finds the least cost path even when the previous least cost path is broken. Popular dynamic routing protocols include distance vector based protocols like Routing Information Protocol (RIP)[3] and link state based protocols like Intermediate System to Intermediate System (IS-IS)[4] and Open Shortest Path First (OSPF)[5].

In this lab, the IS-IS is used as the Interior Gateway Protocol (IGP), which provides faster convergence and larger scalability compared to distance vector based protocols. IS-IS is short for Intermediate System to Intermediate System Routing Protocol. By using this protocol, each router maintains a database which has a map of the whole topology and all routers have the same information. The best path to every destination is computed by all routers. Figure 3.1 shows the design of IS-IS protocol in BT Network. As the figure shows, the network only has level-1 routers for internal routing.

3.1.2 Loopback Addresses and NSAP for Routers

To set up IS-IS, an unique loopback address is needed for each router. An IPv4 address block of 23.0.255.0/24 and IPv6 address block of 2001:2300:ffff::/48 is allocated for loopback addresses.




Figure 3.1: Design of IS-IS Protocol in BT Network.

Following the CLNS addressing convention, each router is then assigned with a NSAP. A NSAP has 3 main components.

According to the convention, the leading Area ID is composed of AFI (49) and Area Address (0001). The System ID followed is set to the IPv4 loopback address of the router. If the loopback address is *ABC.DEF.GHI.JKL*, then System ID should be *ABCD.EFGH.IJKL*. The last main component is N-Selector (NSEL) and set to 00.

The assignment of addresses and NSAP to routers are detailed in Table 3.1.

Router	IPv4 Loopback Address	IPv6 Loopback Address	NSAP
BT-R001	23.0.255.1	2001:2300:ffff:1::	49.0001.0230.0025.5001.00
BT-R002	23.0.255.2	2001:2300:ffff:2::	49.0001.0230.0025.5002.00
BT-R003	23.0.255.3	2001:2300:ffff:3::	49.0001.0230.0025.5003.00

Table 3.1: IP Loopback Addresses and NSAP for Routers in BT Network.

3.1.3 Implementation

IS-IS is set up on Router 1 (BT-R001) using the following commands.

```
1 interface Loopback0
2 ip address 23.0.255.1 255.255.255.255
3 ipv6 address 2001:2300:FFFF:1::/128
4
5 router isis
6 net 49.0001.0230.0025.5001.00
7 is-type level-1
```

Then, IS-IS is turned on on all interfaces to internal routers.

```
1 interface FastEthernet0/0
2 ip router isis
3 ipv6 router isis
4
5 interface FastEthernet0/1
6 ip router isis
7 ipv6 router isis
8
9 interface FastEthernet0/1/0
10 ip router isis
11 ipv6 router isis
```

However, IS-IS routes should not be broadcasted nor received through the loopback interface (Loopback0) while the route to corresponding subnet should be broadcasted to other internal routers. Therefore, the loopback interface should be a passive interface in IS-IS protocol.

```
1 router isis
2 passive-interface Loopback0
```

For Router 2 (BT-R002) and Router 3 (BT-R002), IS-IS is set up similarly using the above commands. The main difference is that interfaces to external routers (VLAN 3 for Router 2 and VLAN 4 & 5 for Router 3) should be passive interfaces as well. Below is the configuration for Router 2.

```
1 interface Loopback0
2 ip address 23.0.255.2 255.255.255.255
3 ipv6 address 2001:2300:FFFF:2::/128
4
5 router isis
6 net 49.0001.0230.0025.5002.00
7 is-type level-1
8 passive-interface Vlan3
9 passive-interface Loopback0
```

```
10
11 interface FastEthernet0/0
12 ip router isis
13 ipv6 router isis
14
15 interface FastEthernet0/1
16 ip router isis
17 ipv6 router isis
18
19 interface Vlan1
20 ip router isis
21 ipv6 router isis
```

Below is the configuration for Router 3.


```
1 interface Loopback0
2 ip address 23.0.255.3 255.255.255.255
3 ipv6 address 2001:2300:FFFF:3::/128
4
5 router isis
6 net 49.0001.0230.0025.5003.00
7 is-type level-1
8 passive-interface Vlan4
9 passive-interface Vlan5
10 passive-interface Loopback0
11
12 interface FastEthernet0/0
13 ip router isis
14 ipv6 router isis
15
16 interface FastEthernet0/1
17 ip router isis
18 ipv6 router isis
19
20 interface Vlan2
21 ip router isis
22 ipv6 router isis
```

3.1.4 Evaluation

Once the IS-IS is set up, use `traceroute` command to check the IPv4 path from Laptop 1 (BT001, IPv4 Address: 23.0.0.2) to Laptop 2 (BT002, IPv4 Address: 23.0.0.18) in the network.

```
1 traceroute 23.0.0.34
```

Figure 3.2 shows the route taken is Laptop 1 (BT001, IPv4 Address: 23.0.0.2) -> Router 1 (BT-R001, IPv4 Address: 23.0.0.1) -> Router 2 (BT-R002, IPv4 Address: 23.0.0.50) -> Laptop 2 (BT002, IPv4 Address: 23.0.0.18), which is a correct route. Routes from Laptop 1 to Laptop 3 as well as from Laptop 2 to Laptop 3 are tested and shown in the figure as well.




```

root@BT001:/home/bt001# traceroute 23.0.0.18
traceroute to 23.0.0.18 (23.0.0.18), 30 hops max, 60 byte packets
 1  23.0.0.33 (23.0.0.33)  1.070 ms  1.643 ms  2.137 ms
 2  23.0.0.53 (23.0.0.53)  1.572 ms  1.930 ms  2.069 ms
 3  23.0.0.18 (23.0.0.18)  1.044 ms  1.023 ms  0.991 ms
root@BT001:/home/bt001# _

```

(a) `traceroute` from Laptop 1 (BT001) to Laptop 2 (BT002)




```

root@BT001:/home/bt001# traceroute 23.0.0.34
traceroute to 23.0.0.34 (23.0.0.34), 30 hops max, 60 byte packets
 1  23.0.0.1 (23.0.0.1)  1.073 ms  1.602 ms  1.840 ms
 2  23.0.0.58 (23.0.0.58)  1.208 ms  1.943 ms  2.264 ms
 3  23.0.0.34 (23.0.0.34)  1.850 ms  1.830 ms  1.809 ms
root@BT001:/home/bt001# _

```

(b) `traceroute` from Laptop 1 (BT001) to Laptop 3 (BT003)



```

root@BT002:/home/bt002# traceroute 23.0.0.34
traceroute to 23.0.0.34 (23.0.0.34), 30 hops max, 60 byte packets
 1  23.0.0.17 (23.0.0.17)  1.006 ms  1.542 ms  1.959 ms
 2  23.0.0.54 (23.0.0.54)  0.931 ms  2.216 ms  2.451 ms
 3  23.0.0.34 (23.0.0.34)  1.856 ms  1.843 ms  1.924 ms
root@BT002:/home/bt002# _

```

(c) `traceroute` from Laptop 2 (BT002) to Laptop 3 (BT003)

Figure 3.2: Tracing IPv4 Routes between Laptops using `traceroute`.

Similar results can be observed for IPv6 routes in Figure 3.3. The IPv6 Route taken from Laptop 1 to Laptop 2 is Laptop 1 (BT001, IPv6 Address: 2001:2300::2) -> Router 1 (BT-R001, IPv6 Address: 2001:2300::1) -> Router 2 (BT-R002, IPv6 Address: 2001:2300:3:2::2) -> Laptop 2 (BT002, IPv6 Address: 2001:2300:0:1::2), which is also a correct route. Routes from Laptop 1 to Laptop 3 as well as from Laptop 2 to Laptop 3 are tested and shown in the figure as well.

```

root@BT001:/home/bt001# traceroute 2001:2300:0:1::2
traceroute to 2001:2300:0:1::2 (2001:2300:0:1::2), 30 hops max, 80 byte packets
 1 2001:2300::1 (2001:2300::1) 1.295 ms 1.656 ms 1.954 ms
 2 2001:2300:0:3::2 (2001:2300:0:3::2) 1.578 ms 1.747 ms 1.858 ms
 3 2001:2300:0:1::2 (2001:2300:0:1::2) 1.088 ms 1.061 ms 1.033 ms
root@BT001:/home/bt001#

```

(a) `traceroute` from Laptop 1 (BT001) to Laptop 2 (BT002)

```

root@BT001:/home/bt001# traceroute 2001:2300:0:2::2
traceroute to 2001:2300:0:2::2 (2001:2300:0:2::2), 30 hops max, 80 byte packets
 1 2001:2300::1 (2001:2300::1) 1.287 ms 1.560 ms 1.804 ms
 2 2001:2300:0:5::2 (2001:2300:0:5::2) 1.500 ms 1.720 ms 1.842 ms
 3 2001:2300:0:2::2 (2001:2300:0:2::2) 1.075 ms 1.061 ms 1.095 ms
root@BT001:/home/bt001#

```

(b) `traceroute` from Laptop 1 (BT001) to Laptop 3 (BT003)

```

bt002@BT002:~$ traceroute 2001:2300:0:2::2
traceroute to 2001:2300:0:2::2 (2001:2300:0:2::2), 30 hops max, 80 byte packets
 1 2001:2300:0:1::1 (2001:2300:0:1::1) 1.446 ms 1.753 ms 2.091 ms
 2 2001:2300:0:4::2 (2001:2300:0:4::2) 1.439 ms 1.632 ms 1.765 ms
 3 2001:2300:0:2::2 (2001:2300:0:2::2) 1.065 ms 1.062 ms 1.036 ms
bt002@BT002:~$

```

(c) `traceroute` from Laptop 2 (BT002) to Laptop 3 (BT003)

Figure 3.3: Tracing IPv6 Routes between Laptops using `traceroute`.

To further evaluate the correctness of our implementation, we check the path from Laptop 1 to Laptop 2 under the condition that the physical connection between Router 1 and Router 2 is broken. The IS-IS protocol on Router 1 should be able to find route to Router 2 through Router 3.

Figure 3.4 shows the IPv4 route taken is Laptop 1 (BT001, IPv4 Address: 23.0.0.2) -> Router 1 (BT-R001, IPv4 Address: 23.0.0.1) -> Router 3 (BT-R003, IPv4 Address: 23.0.0.58) -> Router 2 (BT-R002, IPv4 Address: 23.0.0.53) -> Laptop 2 (BT002, IPv4 Address: 23.0.0.18), which is a correct route.




Figure 3.4: Tracing IPv4 Routes from Laptop 2 (BT002) to Laptop 3 (BT003) when Physical Connection between Router 1 (BT-R001) and Router 2 (BT-R002) is broken.

Figure 3.5 shows the IPv6 route taken is Laptop 1 (BT001, IPv6 Address: 2001:2300::2) -> Router 1 (BT-R001, IPv6 Address: 2001:2300::1) -> Router 3 (BT-R003, IPv6 Address: 2001:2300:0:5::2) -> Router 2 (BT-R002, IPv6 Address: 2001:2300:0:4::1) -> Laptop 2 (BT002, IPv6 Address: 2001:2300:0:1::2), which is also a correct route.



Figure 3.5: Tracing IPv6 Routes from Laptop 2 (BT002) to Laptop 3 (BT003) when Physical Connection between Router 1 (BT-R001) and Router 2 (BT-R002) is broken.

CHAPTER 3. ROUTING PROTOCOLS IN THE NETWORK

On the router's side, `show ip route` and `show ipv6 route` is used to inspect the routes discovered by IS-IS protocol. In Figure 3.6 and 3.7, such routes to other subnets inside BT Network on each router are shown.

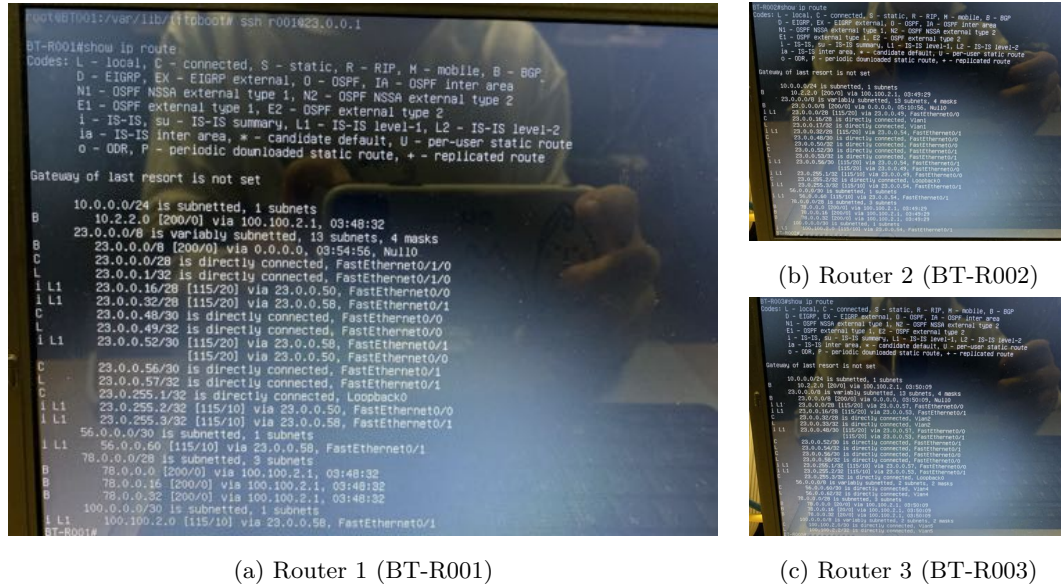


Figure 3.6: IPv4 Routes to Other Subnets on All 3 Routers Respectively using `show ip route`.

CHAPTER 3. ROUTING PROTOCOLS IN THE NETWORK

```
BT-R001#show ipv6 route
IPv6 Routing Table - default - 15 entries
Codes: C - Connected, L - Local, S - Static, U - Per-user Static route
B - BGP, M - MIPv6, R - RIP, I1 - ISIS L1
I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary, O - OSPF
EX - EIGRP external, ND - Neighbor Discovery

B 2001:2300::/32 [200/0]
  via Null0, directly connected
C 2001:2300::/64 [0/0]
  via FastEthernet0/1/0, directly connected
L 2001:2300::1/128 [0/0]
  via FastEthernet0/1/0, receive
I1 2001:2300:0:1::/64 [115/20]
  via FE80::223:EBFF:FEAF:A00, FastEthernet0/0
I1 2001:2300:0:2::/64 [115/20]
  via FE80::223:4FF:FE80:4EEA, FastEthernet0/1
C 2001:2300:0:3::/64 [0/0]
  via FastEthernet0/0, directly connected
L 2001:2300:0:3:1/128 [0/0]
  via FastEthernet0/0, receive
I1 2001:2300:0:4::/64 [115/20]
  via FE80::223:EBFF:FEAF:A00, FastEthernet0/0
  via FE80::223:4FF:FE80:4EEA, FastEthernet0/1
C 2001:2300:0:5::/64 [0/0]
  via FastEthernet0/1, directly connected
L 2001:2300:0:5:1/128 [0/0]
  via FastEthernet0/1, receive
I1 2001:2300:F3:1/128 [115/10]
  via FE80::223:4FF:FE80:4EEA, FastEthernet0/1
LC 2001:2300:FFFF:1:1/128 [0/0]
  via Loopback0, receive
I1 2001:2300:FFFF:12:1/128 [115/10]
  via FE80::223:EBFF:FEAF:A00, FastEthernet0/0
I1 2001:5600:0:6:1/64 [115/10]
  via FE80::223:4FF:FE80:4EEA, FastEthernet0/1
L 2001:2/0 [0/0]
  via Null0, receive
BT-R001#
```

(a) Router 1 (BT-R001)

```
BT-R002#show ipv6 route
IPv6 Routing Table - default - 15 entries
Codes: C - Connected, L - Local, S - Static, U - Per-user Static route
B - BGP, M - MIPv6, R - RIP, I1 - ISIS L1
I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary, O - OSPF
EX - EIGRP external, ND - Neighbor Discovery

B 2001:2300::/32 [200/0]
  via Null0, directly connected
C 2001:2300::/64 [0/0]
  via FastEthernet0/1/0, directly connected
L 2001:2300::1/128 [0/0]
  via FastEthernet0/1/0, receive
I1 2001:2300:0:1::/64 [115/20]
  via FE80::223:EBFF:FEAF:A00, FastEthernet0/0
I1 2001:2300:0:2::/64 [115/20]
  via FE80::223:4FF:FE80:4EEA, FastEthernet0/1
C 2001:2300:0:3::/64 [0/0]
  via FastEthernet0/0, directly connected
L 2001:2300:0:3:1/128 [0/0]
  via FastEthernet0/0, receive
I1 2001:2300:0:4::/64 [115/20]
  via FE80::223:EBFF:FEAF:A00, FastEthernet0/0
  via FE80::223:4FF:FE80:4EEA, FastEthernet0/1
C 2001:2300:0:5::/64 [0/0]
  via FastEthernet0/1, directly connected
L 2001:2300:0:5:1/128 [0/0]
  via FastEthernet0/1, receive
I1 2001:2300:F3:1/128 [115/10]
  via FE80::223:4FF:FE80:4EEA, FastEthernet0/1
LC 2001:2300:FFFF:1:1/128 [0/0]
  via Loopback0, receive
I1 2001:2300:FFFF:12:1/128 [115/10]
  via FE80::223:EBFF:FEAF:A00, FastEthernet0/0
I1 2001:5600:0:6:1/64 [115/10]
  via FE80::223:4FF:FE80:4EEA, FastEthernet0/1
L 2001:2/0 [0/0]
  via Null0, receive
BT-R002#
```

(b) Router 2 (BT-R002)

```
BT-R003#show ipv6 route
IPv6 Routing Table - default - 15 entries
Codes: C - Connected, L - Local, S - Static, U - Per-user Static route
B - BGP, M - MIPv6, R - RIP, I1 - ISIS L1
I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary, O - OSPF
EX - EIGRP external, ND - Neighbor Discovery

B 2001:2300::/32 [200/0]
  via Null0, directly connected
C 2001:2300::/64 [0/0]
  via FastEthernet0/1/0, directly connected
L 2001:2300::1/128 [0/0]
  via FastEthernet0/1/0, receive
I1 2001:2300:0:1::/64 [115/20]
  via FE80::223:EBFF:FEAF:A00, FastEthernet0/0
I1 2001:2300:0:2::/64 [115/20]
  via FE80::223:4FF:FE80:4EEA, FastEthernet0/1
C 2001:2300:0:3::/64 [0/0]
  via FastEthernet0/0, directly connected
L 2001:2300:0:3:1/128 [0/0]
  via FastEthernet0/0, receive
I1 2001:2300:0:4::/64 [115/20]
  via FE80::223:EBFF:FEAF:A00, FastEthernet0/0
  via FE80::223:4FF:FE80:4EEA, FastEthernet0/1
C 2001:2300:0:5::/64 [0/0]
  via FastEthernet0/1, directly connected
L 2001:2300:0:5:1/128 [0/0]
  via FastEthernet0/1, receive
I1 2001:2300:F3:1/128 [115/10]
  via FE80::223:4FF:FE80:4EEA, FastEthernet0/1
LC 2001:2300:FFFF:1:1/128 [0/0]
  via Loopback0, receive
I1 2001:2300:FFFF:12:1/128 [115/10]
  via FE80::223:EBFF:FEAF:A00, FastEthernet0/0
I1 2001:5600:0:6:1/64 [115/10]
  via FE80::223:4FF:FE80:4EEA, FastEthernet0/1
L 2001:2/0 [0/0]
  via Null0, receive
BT-R003#
```

(c) Router 3 (BT-R003)

Figure 3.7: IPv6 Routes to Other Subnets on All 3 Routers Respectively using `show ipv6 route`.

Noticably, when the physical connection between Router 1 and Router 2 is broken, the route from Router 1 to Router 2 goes through Router 3 instead, as evident in Figure 3.8.

3.1.5 Commentary

3.1.5.1 Problem: IS-IS Not Set Up for Laptop-Router Interface

When initially setting up IS-IS on the interfaces, only the interfaces between routers have been turned on. This leads to laptop's failure to reach a router not directly connected. To solve this problem, IS-IS is set up on the interface between a laptop and a router.

3.1.5.2 Alternative Solution to Passive Interfaces

While turning interfaces connected to outside network into passive interfaces does enable other routers in the network to connect to outside networks through those interfaces, it demands extra computing resource for computing routes to external subnets.

Alternatively, one can disable IS-IS on such interfaces and replace the external next-hop for each out-going route in BGP with the internal router to which the next-hop is directly connected to. That prevents Router-Neighbour subnet from participating in IS-IS routing and eliminates the extra computing consumption while preserving the connectivity between internal routers and outside networks.

```

Success rate is 80 percent (4/5), round-trip min/avg/max = 1/2/4 ms
BT-R001#sh ip route
Codes: L - local, C - connected, S - static, R - RIP, M - mobile, B - BGP
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
       E1 - OSPF external type 1, E2 - OSPF external type 2
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
       ia - IS-IS inter area, * - candidate default, U - per-user static route
       o - ODR, P - periodic downloaded static route, + - replicated route

Gateway of last resort is not set

10.0.0.0/24 is subnetted, 1 subnets
B    10.2.2.0 [200/0] via 100.100.2.1, 04:04:39
23.0.0.0/8 is variably subnetted, 12 subnets, 4 masks
B    23.0.0.0/8 [200/0] via 0.0.0.0, 04:11:03, Null0
C    23.0.0.0/28 is directly connected, FastEthernet0/1/0
L    23.0.0.1/32 is directly connected, FastEthernet0/1/0
i L1 23.0.0.16/28 [115/30] via 23.0.0.58, FastEthernet0/1
i L1 23.0.0.32/28 [115/20] via 23.0.0.58, FastEthernet0/1
i L1 23.0.0.48/30 [115/30] via 23.0.0.58, FastEthernet0/1
i L1 23.0.0.52/30 [115/20] via 23.0.0.58, FastEthernet0/1
C    23.0.0.56/30 is directly connected, FastEthernet0/1
L    23.0.0.57/32 is directly connected, FastEthernet0/1
C    23.0.255.1/32 is directly connected, Loopback0
i L1 23.0.255.2/32 [115/20] via 23.0.0.58, FastEthernet0/1
i L1 23.0.255.3/32 [115/10] via 23.0.0.58, FastEthernet0/1
56.0.0.0/30 is subnetted, 1 subnets
i L1 56.0.0.60 [115/10] via 23.0.0.58, FastEthernet0/1
78.0.0.0/28 is subnetted, 3 subnets
B    78.0.0.0 [200/0] via 100.100.2.1, 04:04:39
B    78.0.0.16 [200/0] via 100.100.2.1, 04:04:39
B    78.0.0.32 [200/0] via 100.100.2.1, 04:04:39
100.0.0.0/30 is subnetted, 1 subnets
i L1 100.100.2.0 [115/10] via 23.0.0.58, FastEthernet0/1
BT-R001#

```

(a) IPv4 Routes to Other Subnets on Router 1 using show ip route.

```

BT-R001#sh ipv6 route
IPv6 Routing Table - default - 13 entries
Codes: C - Connected, L - Local, S - Static, U - Per-user Static route
       B - BGP, M - MIPv6, R - RIP, I1 - ISIS L1
       I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary, D - EIGRP
       EX - EIGRP external, ND - Neighbor Discovery

B    2001:2300::/32 [200/0]
    via Null0, directly connected
C    2001:2300::/64 [0/0]
    via FastEthernet0/1/0, directly connected
L    2001:2300::1/128 [0/0]
    via FastEthernet0/1/0, receive
I1   2001:2300:0:1::/64 [115/30]
    via FE80::223:4FF:FE80:4EEA, FastEthernet0/1
I1   2001:2300:0:2::/64 [115/20]
    via FE80::223:4FF:FE80:4EEA, FastEthernet0/1
I1   2001:2300:0:4::/64 [115/20]
    via FE80::223:4FF:FE80:4EEA, FastEthernet0/1
C    2001:2300:0:5::/64 [0/0]
    via FastEthernet0/1, directly connected
L    2001:2300:0:5::1/128 [0/0]
    via FastEthernet0/1, receive
I1   2001:2300:F:3::/128 [115/10]
    via FE80::223:4FF:FE80:4EEA, FastEthernet0/1
LC   2001:2300:FFFF:1::/128 [0/0]
    via Loopback0, receive
I1   2001:2300:FFFF:2::/128 [115/20]
    via FE80::223:4FF:FE80:4EEA, FastEthernet0/1
I1   2001:5600:0:6::/64 [115/10]
    via FE80::223:4FF:FE80:4EEA, FastEthernet0/1
L    FF00::/8 [0/0]
    via Null0, receive
BT-R001#

```

(b) IPv6 Routes to Other Subnets on Router 1 using show ip route.

Figure 3.8: IP Routes to Other Subnets on Router 1 when Physical Connection between Router 1 (BT-R001) and Router 2 (BT-R002) is broken.

3.2 Inter-domain Routing Protocol: BGP

3.2.1 Design

For inter-domain routing, Border Gateway Protocol (BGP)[6] is applied in BT Network since it's the one and only External Gateway Protocol (EGP) in today's global Internet. BGP provides scalability to large networks, clear definitions of administrative boudnary as well as flexiable policy control, which allow business relaitonships with neighbouring ISPs to be expressed in terms of routing policies.

Figure 3.9 shows the design of BGP protocol in our network. The AS Number (ASN) of BT Network is 2030 while ASN of Central, Virgin, DT are 42, 5060, 3040 respectively.




Figure 3.9: Design of BGP Protocol in BT Network.

Router 2 (BT-R002) and Router 3 (BT-R003) act as External BGP routers since they are directly connected to neighbouring ISPs. They receive routes announced by neighbouring ISPs' routers and announce routes originated from BT Network. Router 1 (BT-R001) acts

as an Internal BGP (IBGP) router only since it is not directly connected to any neighbouring ISP and only forwards and receives routes announced by EGBP routers.

3.2.2 Routing Policies

To enact business relationships with neighbouring ISPs, proper routing policies should be implemented for each different relationship in BGP protocol. Specifically, for customer ISPs, all routes announced by such ISP should be accepted and all routes received should be announced to it as well. This allows customers to connect to BT Network as well as connect to other networks through BT Network.

For non-customer neighbouring ISPs (peers and providers), all routes announced by such ISP should be accepted in order for BT Network to connect to it. Meanwhile, only routes originated from either BT Network or BT's customers are announced to such ISP.

3.2.3 Implementation

On Router 1 (BT-R001, IPv4 Loopback Address: 23.0.255.1, IPv6 Loopback Address: 2001:2300:FFFF:1::), both Router 2 (BT-R002, IPv4 Loopback Address: 23.0.255.2, IPv6 Loopback Address: 2001:2300:FFFF:2::) and Router 3 (BT-R003, IPv4 Loopback Address: 23.0.255.3, IPv6 Loopback Address: 2001:2300:FFFF:3::) are taken as neighbouring routers in the same AS. In addition, the source of BGP messages are set to be the loopback address of Router 1 to prevent physical disconnection to the 2 routers. Router 1 announces the subnet BT-R001 - BT001 (IPv4: 23.0.0.0/28, IPv6: 2001:2300:0:0::/64) to other routers.

```
1 router bgp 2030
2 network 23.0.0.0 mask 255.255.255.240
3 neighbor 23.0.255.2 remote-as 2030
4 neighbor 23.0.255.2 update-source Loopback0
5 neighbor 23.0.255.3 remote-as 2030
6 neighbor 23.0.255.3 update-source Loopback0
7 neighbor 2001:2300:FFFF:2:: remote-as 2030
8 neighbor 2001:2300:FFFF:2:: update-source Loopback0
9 neighbor 2001:2300:FFFF:3:: remote-as 2030
10 neighbor 2001:2300:FFFF:3:: update-source Loopback0
11
12 address-family ipv6
13 network 2001:2300::/64
14 neighbor 2001:2300:FFFF:2:: activate
15 neighbor 2001:2300:FFFF:3:: activate
```

On Router 2 (BT-R002), both Router 1 and Router 3 are taken as neighbouring routers in the same AS while the router from DT is taken as router from AS 3040. Router 2 announces

the subnet BT-R002 - BT002 (IPv4: 23.0.0.16/28, IPv6: 2001:2300:0:1::/64) to other routers. Since Router 2 is directly connected to customer DT Network, it applies no filter on inbound and outbound routes.

```

1 router bgp 2030
2 network 23.0.0.16 mask 255.255.255.240
3 neighbor 23.0.0.62 remote-as 3040
4 neighbor 23.0.255.1 remote-as 2030
5 neighbor 23.0.255.1 update-source Loopback0
6 neighbor 23.0.255.3 remote-as 2030
7 neighbor 23.0.255.3 update-source Loopback0
8 neighbor 2001:2300:0:6::2 remote-as 3040
9 neighbor 2001:2300:FFFF:1:: remote-as 2030
10 neighbor 2001:2300:FFFF:1:: update-source Loopback0
11 neighbor 2001:2300:FFFF:3:: remote-as 2030
12 neighbor 2001:2300:FFFF:3:: update-source Loopback0
13
14 address-family ipv6
15 network 2001:2300:0:1::/64
16 neighbor 2001:2300:0:6::2 activate
17 neighbor 2001:2300:FFFF:1:: activate
18 neighbor 2001:2300:FFFF:3:: activate

```

On Router 3 (BT-R003), both Router 1 and Router 2 are taken as neighbouring routers in the same AS while the routers from Central and Virgin are taken as router from AS 42 and AS 5060 respectively. Router 3 announces the subnet BT-R003 - BT003 (IPv4: 23.0.0.32/28, IPv6: 2001:2300:0:2::/64) to other routers.

```

1 router bgp 2030
2 network 23.0.0.32 mask 255.255.255.240
3 neighbor 23.0.255.1 remote-as 2030
4 neighbor 23.0.255.1 update-source Loopback0
5 neighbor 23.0.255.2 remote-as 2030
6 neighbor 23.0.255.2 update-source Loopback0
7 neighbor 2001:2300:FFFF:1:: remote-as 2030
8 neighbor 2001:2300:FFFF:1:: update-source Loopback0
9 neighbor 2001:2300:FFFF:2:: remote-as 2030
10 neighbor 2001:2300:FFFF:2:: update-source Loopback0
11 neighbor 2001:5600:0:6::1 remote-as 5060
12 neighbor 56.0.0.61 remote-as 5060
13 neighbor 100.100.2.1 remote-as 42
14
15 address-family ipv6
16 network 2001:2300:0:2::/64
17 aggregate-address 2001:2300::/32 summary-only
18 neighbor 2001:2300:FFFF:1:: activate
19 neighbor 2001:2300:FFFF:2:: activate

```


Since Router 3 is directly connected to non-customer ISPs, it applies a filter on outbound routes, which denies all routes that pass through either Central Network (ASN: 42) or Virgin Network (ASN: 5060).

```
1 ip as-path access-list 1 deny _42_  
2 ip as-path access-list 1 deny _5060_  
3 ip as-path access-list 1 permit .  
4 router bgp 2030  
5 neighbor 56.0.0.61 filter-list 1 out  
6 neighbor 100.100.2.1 filter-list 1 out
```

Address aggregation for 23.0.0.0/8 and 2001:2300::/32 is set up on all 3 routers, which aggregates routes destined for all addresses inside BT Network range into a single route.

```
1 aggregate-address 23.0.0.0 255.0.0.0 summary-only  
2 address-family ipv6  
3 aggregate-address 2001:2300::/32 summary-only
```

3.2.4 Evaluation

3.2.4.1 BGP Routes

Routes collected through BGP protocol on all 3 routers are shown in Figure 3.10 and 3.11 using commands `show bgp` and `show bgp ipv6`. Routes to DT Network (IPv4: 34.0.0.0/8, IPv6: 2001:3400::/32), Virgin Network (IPv4: 56.0.0.0/8, IPv6: 2001:5600::/32), Central Network (IPv4: 10.2.2.0/24) and other networks can be observed in the figure.


```

root@BT001:/home/bt001# sudo su
root@BT001:/home/bt001# ssh r001@23.0.0.1

BT-R001#sh bgp ipv6
BGP table version is 22, local router ID is 23.0.255.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

   Network        Next Hop           Metric LocPrf Weight Path
s> 2001:2300::/64  ::                0         32768 i
* 12001:2300::/32  2001:2300:FFFF:3:: 0         100      0 i
* 1                2001:2300:FFFF:2:: 0         100      0 i
*>                ::                0         32768 i
*> 12001:5600::/64  2001:5600:0:6::1   20        100      0 5060 i
*> 12001:5600:0:1::/64  2001:5600:0:6::1   0         100      0 5060 i
*> 12001:5600:0:2::/64  2001:5600:0:6::1   20        100      0 5060 i
*> 12001:6700:0:1::/64  2001:5600:0:6::1   0         100      0 5060 6070 i
*> 12001:6700:0:2::/64  2001:5600:0:6::1   0         100      0 5060 6070 i
*> 12001:6700:0:3::/64  2001:5600:0:6::1   0         100      0 5060 6070 i

% NOTE: This command is deprecated. Please use 'show bgp ipv6 unicast'
BT-R001#


```

(a) Router 1 (BT-R001)

Figure 3.11: IPv6 Routes Collected through BGP Protocols on Routers using `show bgp ipv6`.

3.2.4.2 Connectivity to Provider Central Network

The connectivity to provider Central Network using BGP protocol is tested and evaluated by tracing routes to IP addresses 10.2.2.1 on Laptop 1 (BT001). As shown in Figure 3.12, connection to Central Network is successfully established through BGP routes.



(a) 10.2.2.1

Figure 3.12: Tracing IPv4 Routes to Central Network on Laptop 1 (BT001) using traceroute.

3.2.4.3 Connectivity to Customer DT Network

The connectivity to customer DT Network using BGP protocol is tested and evaluated by tracing routes to IP addresses 34.0.0.2, 34.0.0.18 and 34.0.0.34 on Laptop 1 (BT001). As shown in Figure 3.13, connection to DT Network is successfully established through BGP routes.

```

root@BT001:/home/bt001# traceroute 34.0.0.2
traceroute to 34.0.0.2 (34.0.0.2), 30 hops max, 60 byte packets
 1 23.0.0.1 (23.0.0.1)  1.225 ms  1.424 ms  1.736 ms
 2 23.0.0.50 (23.0.0.50)  1.377 ms  1.668 ms  2.059 ms
 3 23.0.0.62 (23.0.0.62)  1.779 ms  2.294 ms  2.740 ms
 4 34.0.0.53 (34.0.0.53)  1.963 ms  2.218 ms  2.627 ms
 5 34.0.0.2 (34.0.0.2)  1.631 ms  1.612 ms  2.556 ms
root@BT001:/home/bt001#

```

(a) 34.0.0.2

```

root@BT001:/home/bt001# traceroute 34.0.0.18
traceroute to 34.0.0.18 (34.0.0.18), 30 hops max, 60 byte packets
 1 23.0.0.1 (23.0.0.1)  1.245 ms  1.441 ms  1.861 ms
 2 23.0.0.50 (23.0.0.50)  1.406 ms  1.810 ms  2.216 ms
 3 23.0.0.62 (23.0.0.62)  2.020 ms  2.367 ms  2.827 ms
 4 34.0.0.18 (34.0.0.18)  1.536 ms  1.487 ms  1.467 ms
root@BT001:/home/bt001#

```

(b) 34.0.0.18

```

root@BT001:/home/bt001# traceroute 34.0.0.34
traceroute to 34.0.0.34 (34.0.0.34), 30 hops max, 60 byte packets
 1 23.0.0.1 (23.0.0.1)  1.231 ms  1.424 ms  1.733 ms
 2 23.0.0.50 (23.0.0.50)  1.376 ms  1.666 ms  2.049 ms
 3 23.0.0.62 (23.0.0.62)  1.870 ms  2.317 ms  2.909 ms
 4 34.0.0.57 (34.0.0.57)  2.109 ms  2.333 ms  2.652 ms
 5 34.0.0.34 (34.0.0.34)  1.658 ms  1.642 ms  2.557 ms
root@BT001:/home/bt001# _

```

(c) 34.0.0.34

Figure 3.13: Tracing IPv4 Routes to DT Network on Laptop 1 (BT001) using `traceroute`.

3.2.4.4 Connectivity to Peer Virgin Network

The connectivity to peer Virgin Network using BGP protocol is tested and evaluated by tracing routes to IP addresses 56.0.0.2, 56.0.0.18 and 56.0.0.34 on Laptop 1 (BT001). As shown in Figure 3.15, connection to Virgin Network is successfully established through BGP routes.

```
root@BT001:/etc/bind# traceroute 56.0.0.2
traceroute to 56.0.0.2 (56.0.0.2), 30 hops max, 60 byte packets
 1  23.0.0.1 (23.0.0.1)  1.206 ms  1.494 ms  1.804 ms
 2  23.0.0.58 (23.0.0.58)  1.425 ms  1.759 ms  2.273 ms
 3  56.0.0.61 (56.0.0.61)  1.873 ms  2.329 ms  2.762 ms
 4  56.0.0.49 (56.0.0.49)  1.948 ms  2.237 ms  2.442 ms
 5  56.0.0.2 (56.0.0.2)  1.680 ms  1.663 ms  1.641 ms
root@BT001:/etc/bind#
```

(a) 56.0.0.2

```
root@BT001:/etc/bind# traceroute 56.0.0.18
traceroute to 56.0.0.18 (56.0.0.18), 30 hops max, 60 byte packets
 1  23.0.0.1 (23.0.0.1)  1.201 ms  1.489 ms  1.902 ms
 2  23.0.0.58 (23.0.0.58)  1.421 ms  1.968 ms  2.287 ms
 3  56.0.0.61 (56.0.0.61)  1.919 ms  2.252 ms  2.563 ms
 4  56.0.0.18 (56.0.0.18)  1.451 ms  1.431 ms  1.410 ms
root@BT001:/etc/bind#
```

(b) 56.0.0.18


```
root@BT001:/etc/bind# traceroute 56.0.0.34
traceroute to 56.0.0.34 (56.0.0.34), 30 hops max, 60 byte packets
 1  23.0.0.1 (23.0.0.1)  1.070 ms  1.610 ms  1.827 ms
 2  23.0.0.58 (23.0.0.58)  1.182 ms  1.778 ms  2.073 ms
 3  56.0.0.61 (56.0.0.61)  2.478 ms  2.921 ms  3.418 ms
 4  56.0.0.58 (56.0.0.58)  2.860 ms  3.147 ms  3.346 ms
 5  56.0.0.34 (56.0.0.34)  2.494 ms  2.477 ms  2.460 ms
root@BT001:/etc/bind#
```

(c) 56.0.0.34

Figure 3.14: Tracing IPv4 Routes to Virgin Network on Laptop 1 (BT001) using traceroute.

3.2.4.5 Connectivity to Other Networks

The connectivity to other networks using BGP protocol is tested and evaluated by tracing routes to IP addresses 78.0.0.2 (Sonara Network) and 89.0.0.18 (NTT Network) on Laptop 1 (BT001). As shown in Figure 3.15, connection to Virgin Network is successfully established through BGP routes.



```
root@BT001:/etc/bind# traceroute 78.0.0.2
traceroute to 78.0.0.2 (78.0.0.2), 30 hops max, 60 byte packets
 1 23.0.0.1 (23.0.0.1)  1.233 ms  1.429 ms  1.725 ms
 2 23.0.0.50 (23.0.0.50) 1.379 ms  1.643 ms  1.894 ms
 3 23.0.0.62 (23.0.0.62) 1.856 ms  2.144 ms  2.683 ms
 4 34.0.0.66 (34.0.0.66) 2.074 ms  2.321 ms  2.608 ms
 5 78.1.0.1 (78.1.0.1) 2.255 ms * *
root@BT001:/etc/bind# _
```

(a) 78.0.0.2



```
root@BT001:/etc/bind# traceroute 89.0.0.18
traceroute to 89.0.0.18 (89.0.0.18), 30 hops max, 60 byte packets
 1 23.0.0.1 (23.0.0.1)  1.236 ms  1.431 ms  1.795 ms
 2 23.0.0.50 (23.0.0.50) 1.395 ms  1.744 ms  2.152 ms
 3 23.0.0.62 (23.0.0.62) 1.904 ms  2.383 ms  2.667 ms
 4 34.0.0.66 (34.0.0.66) 429.823 ms 429.804 ms 430.178 ms
 5 89.0.0.18 (89.0.0.18) 1.709 ms  1.702 ms  1.678 ms
root@BT001:/etc/bind#
```

(b) 89.0.0.18

Figure 3.15: Tracing IPv4 Routes to Other Networks on Laptop 1 (BT001) using traceroute.

3.2.4.6 Connectivity to Virgin when Direct Physical Connection Is Down

In addition, the connectivity to peer Virgin Network under the unfortunate condition that the direct physical connection is down is also tested. Traced routes to IP addresses 56.0.0.2, 56.0.0.18 and 56.0.0.34 on Laptop 1 (BT001) are shown in Figure 3.16.

Alternative connection to Virgin Network is successfully established through Central Network (ASN: 42).

3.2.5 Commentary

3.2.5.1 Problem: Filter List Not Working for Self-Originated Routes

The initial filter list for outbound routes on Router 3 only permits routes originated from BT Network (ASN 2030) and customer DT Network (ASN 3040).

```
1 ip as-path access-list 1 permit _2030$
2 ip as-path access-list 1 permit _3040$
```

However, the filter list blocks all routes except those originated from customer DT Network to be announced. To solve this problem, a filter list where all routes except those go through provider Central Network (ASN 42) and peer Virgin Network (ASN 5060) are allowed. The new list should have the same effects as the previous list and indeed works as intended.

```
1 ip as-path access-list 1 deny _42_
2 ip as-path access-list 1 deny _5060_
3 ip as-path access-list 1 permit .*
```



```

root@BT001:/etc/bind# traceroute 56.0.0.2
traceroute to 56.0.0.2 (56.0.0.2), 30 hops max, 60 byte packets
 1  23.0.0.1 (23.0.0.1)  1.230 ms  1.582 ms  1.844 ms
 2  23.0.0.58 (23.0.0.58)  1.662 ms  2.113 ms  2.464 ms
 3  100.100.2.1 (100.100.2.1)  1.271 ms  1.692 ms  1.822 ms
 4  100.100.5.2 (100.100.5.2)  1.785 ms  2.188 ms  2.580 ms
 5  56.0.0.53 (56.0.0.53)  2.138 ms  2.510 ms  2.792 ms
 6  56.0.0.2 (56.0.0.2)  2.059 ms  0.875 ms  0.830 ms
root@BT001:/etc/bind# _

```

(a) 56.0.0.2

```

root@BT001:/etc/bind# traceroute 56.0.0.18
traceroute to 56.0.0.18 (56.0.0.18), 30 hops max, 60 byte packets
 1  23.0.0.1 (23.0.0.1)  1.056 ms  1.571 ms  1.817 ms
 2  23.0.0.58 (23.0.0.58)  1.185 ms  1.769 ms  2.139 ms
 3  100.100.2.1 (100.100.2.1)  1.909 ms  4.772 ms  5.014 ms
 4  100.100.5.2 (100.100.5.2)  2.369 ms  2.834 ms  3.184 ms
 5  56.0.0.57 (56.0.0.57)  2.789 ms  3.117 ms  3.415 ms
 6  56.0.0.18 (56.0.0.18)  2.430 ms  0.935 ms  0.829 ms
root@BT001:/etc/bind#

```

(b) 56.0.0.18

```

root@BT001:/etc/bind# traceroute 56.0.0.34
traceroute to 56.0.0.34 (56.0.0.34), 30 hops max, 60 byte packets
 1  23.0.0.1 (23.0.0.1)  1.230 ms  1.594 ms  1.862 ms
 2  23.0.0.58 (23.0.0.58)  1.656 ms  2.073 ms  2.484 ms
 3  100.100.2.1 (100.100.2.1)  1.274 ms  1.965 ms  2.148 ms
 4  100.100.5.2 (100.100.5.2)  1.742 ms  2.081 ms  2.325 ms
 5  56.0.0.34 (56.0.0.34)  1.388 ms  1.370 ms  1.611 ms
root@BT001:/etc/bind# _

```

(c) 56.0.0.34

Figure 3.16: Tracing IPv4 Routes to Virgin Network on Laptop 1 (BT001) using `traceroute` When Direct Physical Connection Is Broken .

Chapter 4

Applications in the Network

4.1 Secure Remote Access to Routers through SSH

4.1.1 Design

Accessing the routers through the physical "console" port is inconvenient and dangerous. Thus, remote access through Secure Shell (SSH) protocol[7] to routers is needed.

In BT network, remote SSH access is enabled on all 3 routers. Seperate combinations of username and password on each router are used to ensure the independence of security of each router.

In addition, SSH public key authentication is set up on Laptop 1 (BT001), which allows the root user on the laptop to login in to all routers without entering passwords.

4.1.2 Implementation

We first set up Remote SSH access was first set up as instructed in Reference Guide on all 3 routers. Below is the configuration commands for Router 1 (BT-R001).

```
1 hostname BT-R001
2 ip domain name bt.lboro
3 username r001 priv 15 secret <secret>
4 line vty 0 4
5 transport input ssh telnet
6 login local
7
8 ip ssh version 2
9 crypto key generate rsa general-keys
10 ip ssh dh min size 4096
```

We then generate a pair of public and private keys on Laptop 1 (BT001).

```
1 ssh-keygen
```

After that, the pair of keys is written into files `~/.ssh/id_rsa` and `~/.ssh/id_rsa.pub`. We use the generated public key (`id_rsa.pub`) to set up SSH public key authentication on all 3 routers.


```
1 ip ssh pubkey-chain
2 username r001
3 key-string
```

4.1.3 Evaluation


Once remote SSH access is set up on 3 routers, one should be able to access them on Laptop 1 (BT001) without entering the password using the following commands.

```
1 # access Router 1
2 ssh r001@23.0.0.1
3 # access Router 2
4 ssh r002@23.0.0.50
5 # access Router 3
6 ssh r003@23.0.0.33
```


Screenshots of successful remote access to all 3 routers are shown in Figure 4.1.



(a) Router 1 (BT-R001)



(b) Router 2 (BT-R002)



(c) Router 3 (BT-R003)

Figure 4.1: Successful remote SSH access to all 3 routers from Laptop 1 (BT001).

4.1.4 Commentary

4.1.4.1 Problem: Maximum Limit of Characters per Line

When we tried to set up SSH public key authentication on routers, we failed at our initial attempt. It turned out that Cisco router has maximum limit of characters for each command line. Thus, a public key in a single long line was not accepted by the router.

To solve this problem, `fold` command is used to split the public key into multiple lines before re-uploading the key and SSH public key authentication was successfully set up on the router.

4.2 Domain Name System Service

4.2.1 Design

Domain Name System (DNS)[8] translates IP addresses to domain names and vice versa. In contrast to hard-to-remember IP addresses, short and meaningful domain names (eg. `lboro.ac.uk` for Loughborough University in UK) are more convenient for Internet users. Additionally, the service providers can change the IP addresses of servers without re-notifying their customers.

In BT Network, there are two DNS servers, one primary master server and the other secondary. Primary master server is deployed at Laptop 3 (BT003) and secondary master server is at Laptop 1 (BT001). The rationale is that when the primary becomes unavailable, the secondary can be the backup domain name server.

Both are authoritative of domain `bt.lboro`. **Each laptop in the network has a corresponding domain name. For example, the domain name of Laptop 1 (BT001) is `bt001.bt.lboro`. In addition, both A and AAAA records of `bt.lboro` point to Laptop 1 while the MX record points to Laptop 3.**

The two DNS servers are also connected to the central DNS server, which is authoritative of domain `lboro`.

4.2.2 Implementation

Install `bind9` package and `dnsutils` package on Laptop 1 (BT001) and Laptop 3 (BT001) using the following commands. All DNS configurations are all stored in folder `/etc/bind`.

```
1 sudo apt-get install bind9
2 sudo apt-get install dnsutils
```

On the primary DNS server, forward unknown DNS requests to central DNS server by adding this line to file `named.conf.options`.

```
1 forwarders { 10.2.2.1; };
```

The following lines are added to file `/etc/bind/named.conf.local`.

The zone section defines the type of the DNS server and it is stored in a file mentioned in the 'file' field. The 'allow-transfer' field defines a match list which has IP addresses that are allowed to do transfer and copy operations to the zone information with the server. The 'allow-notify' field defines an IP addresses match list that is allowed to notify this server and implicitly update the zone. In this case, both fields should be the IP address of the secondary DNS server (Laptop 1).

The first zone is a forward zone which translate the domain name to IP address. The zone name of it is the selected domain name `bt.lboro`. The second zone is the reverse zone which translate the IP address to domain name. The zone name of it should be the fixed IP prefix part and host part. In this case, the name is `23.in-addr.arpa`.

```
1 zone bt.lboro {
2     type master;
3     file /etc/bind/db.bt.lboro ;
4     allow-transfer { 23.0.0.2; };
5     also-notify { 23.0.0.2; };
6 };
7
8 zone 23.in-addr.arpa {
9     type master;
10    file /etc/bind/db.23 ;
11    allow-transfer { 23.0.0.2; };
12    also-notify { 23.0.0.2; };
13 };
```

Then, the files related to forward zone and reverse zone are edited. The file `db.bt.lboro` whose contents are shown in Listing 4.1 defines the forward DNS configuration.

Listing 4.1: Contents of Forward DNS Configuration File Located at `/etc/bind/db.bt.lboro` on Primary DNS Server (Laptop 3).

```
1 ;
2 ; BIND data file for BT.LBORO
3 ;
4 $TTL 604800
5 @ IN SOA bt.lboro. root.bt.lboro. (
6     13 ; Serial
7     604800 ; Refresh
8     86400 ; Retry
9     2419200 ; Expire
10    604800 ) ; Negative Cache TTL
11 ;
12 @ IN NS ns.bt.lboro.
13 @ IN NS ns2.bt.lboro.
14 www IN CNAME bt.lboro.
15 ns IN A 23.0.0.34
16 ns IN AAAA 2001:2300:0:2::2
17 ns2 IN A 23.0.0.2
18 ns2 IN AAAA 2001:2300:0:0::2
19 @ IN A 23.0.0.2
20 @ IN AAAA 2001:2300:0:0::2
21 @ IN MX 10 bt003.bt.lboro.
22 bt001 IN A 23.0.0.2
```

CHAPTER 4. APPLICATIONS IN THE NETWORK

```
23 bt001 IN AAAA 2001:2300:0:0::2
24 bt002 IN A 23.0.0.18
25 bt002 IN AAAA 2001:2300:0:1::2
26 bt003 IN A 23.0.0.34
27 bt003 IN AAAA 2001:2300:0:2::2
```

Serial number is the version number of this file and it should be increased after changing the file. For each line of records, the prefix (eg. `bt001`), domain type (eg. `A`) and value (eg. `23.0.0.2`) are specified.

Listing 4.2: Contents of Reverse DNS Configuration File Located at `/etc/bind/db.23` on Primary DNS Server (Laptop 3).

```
1 ;
2 ; BIND reverse data file for 23.xxx.xxx.xxx net
3 ;
4 $TTL 604800
5 @ IN SOA ns.bt.lboro. root.bt.lboro. (
6     7 ; Serial
7     604800 ; Refresh
8     86400 ; Retry
9     2419200 ; Expire
10    604800 ) ; Negative Cache TTL
11 ;
12 @ IN NS ns.
13 2.0.0 IN PTR bt.lboro.
14 34.0.0 IN PTR ns.bt.lboro.
15 2.0.0 IN PTR ns2.bt.lboro.
16 2.0.0 IN PTR bt001.bt.lboro.
17 18.0.0 IN PTR bt002.bt.lboro.
18 34.0.0 IN PTR bt003.bt.lboro.
```

Listing 4.2 shows the configuration for the reverse zone. The serial number needs to be increased on each change as well. For each line of records, the first column of match list should be reverse host part of IP address. For example, the IP address of Laptop 1 (BT001) is `23.0.0.2` and the host part of it is `0.0.2`. Therefore, the reverse host part of it should be `2.0.0`.

For each A record in forward configurations, a PTR records should be added to the reverse configuration.

When all configurations for DNS server are completed, use following commands to restart `bind9` service to take effect.

```
1 service bind9 restart
```

As for secondary master, Laptop 1 (BT001) only needs to add the following lines to the file `/etc/bind/named.conf.local`, which specify both forward and reverse DNS zones and

Laptop 3 as their master server.

```
1 zone    bt.lboro {
2         type slave;
3         file  db.bt.lboro ;
4         masters { 23.0.0.34; };
5 };
6
7 zone    23.in-addr.arpa {
8         type slave;
9         file  db.23 ;
10        masters { 23.0.0.34; }
11 };
```

Restart `bind9` service for configurations to take effect.

```
1 service bind9 restart
```

Now, Laptop 1 should be able to receive DNS records from the master server and act as a secondary DNS server.

On client's side, Laptop 3 and 2 are set as nameservers in file `/etc/network/interfaces` on all 3 laptops.

```
1 dns-nameservers 23.0.0.34 23.0.0.2
```

Full configurations for both primary and secondary DNS server are detailed in Appendix D.

4.2.3 Evaluation

The command `dig` is used to query DNS records on Laptop 2 for evaluation. In Figure 4.2, DNS records of `bt.lboro` and `www.bt.lboro` are "digged" and correct answers are returned by DNS servers.

In addition, DNS records of `bt001.bt.lboro`, `bt002.bt.lboro` and `bt003.bt.lboro` are "digged" on Laptop 2 as well. In Figure 4.3, correct answers are returned by DNS servers.

In terms of IPv6 addresses, DNS AAAA records of `bt001.bt.lboro`, `bt002.bt.lboro` and `bt003.bt.lboro` are "digged" on Laptop 2 as well. In Figure 4.4, correct answers are returned by DNS servers.

For reverse DNS, records of Laptop 1 (23.0.0.2), Laptop 2 (23.0.0.18) and Laptop 3 (23.0.0.34) are "digged" on Laptop 2 as well. In Figure 4.5, correct answers are returned by DNS servers.

In terms of connectivity, all three laptops are able to directly `ping` the domain name `bt.lboro`, whose DNS A record points to Laptop 1 (BT001, IPv4 Address: 23.0.0.2), as shown in Figure 4.6.

CHAPTER 4. APPLICATIONS IN THE NETWORK




Figure 4.2: Querying DNS Records for `bt.lboro` and `www.bt.lboro` on Laptop 2.




Figure 4.3: Querying DNS Records for sub-domains of `bt.lboro` on Laptop 2.

CHAPTER 4. APPLICATIONS IN THE NETWORK

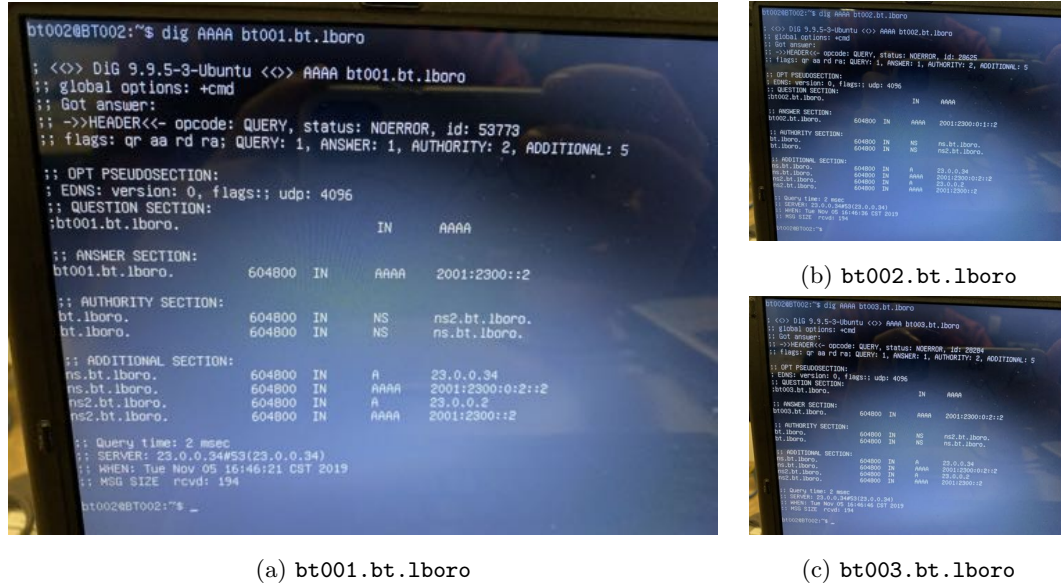


Figure 4.4: Querying DNS AAAA Records for sub-domains of bt.lboro on Laptop 2.



Figure 4.5: Querying Reverse DNS Records for All 3 Laptops on Laptop 2.

```
? ^CInterrupt
? exit
bt003@BT003:~$ ping bt.lboro
PING bt.lboro (23.0.0.2) 56(84) bytes of data:
64 bytes from 23.0.0.2: icmp_seq=1 ttl=62 time=0.845 ms
64 bytes from 23.0.0.2: icmp_seq=2 ttl=62 time=0.884 ms
64 bytes from 23.0.0.2: icmp_seq=3 ttl=62 time=0.862 ms
64 bytes from 23.0.0.2: icmp_seq=4 ttl=62 time=0.868 ms
64 bytes from 23.0.0.2: icmp_seq=5 ttl=62 time=0.878 ms
64 bytes from 23.0.0.2: icmp_seq=6 ttl=62 time=0.763 ms
64 bytes from 23.0.0.2: icmp_seq=7 ttl=62 time=0.874 ms
64 bytes from 23.0.0.2: icmp_seq=8 ttl=62 time=0.882 ms
64 bytes from 23.0.0.2: icmp_seq=9 ttl=62 time=0.879 ms
64 bytes from 23.0.0.2: icmp_seq=10 ttl=62 time=0.862 ms
64 bytes from 23.0.0.2: icmp_seq=11 ttl=62 time=0.880 ms
^C
--- bt.lboro ping statistics ---
11 packets transmitted, 11 received, 0% packet loss, time 10010ms
rtt min/avg/max/mdev = 0.763/0.861/0.884/0.045 ms
bt003@BT003:~$
```

(a) Laptop 1 (BT001)

```
bt002@BT002:~$ ping bt.lboro
PING bt.lboro (23.0.0.2) 56(84) bytes of data:
64 bytes from 23.0.0.2: icmp_seq=1 ttl=62 time=0.745 ms
64 bytes from 23.0.0.2: icmp_seq=2 ttl=62 time=0.802 ms
64 bytes from 23.0.0.2: icmp_seq=3 ttl=62 time=0.872 ms
64 bytes from 23.0.0.2: icmp_seq=4 ttl=62 time=0.877 ms
64 bytes from 23.0.0.2: icmp_seq=5 ttl=62 time=0.883 ms
64 bytes from 23.0.0.2: icmp_seq=6 ttl=62 time=0.887 ms
64 bytes from 23.0.0.2: icmp_seq=7 ttl=62 time=0.888 ms
64 bytes from 23.0.0.2: icmp_seq=8 ttl=62 time=0.890 ms
64 bytes from 23.0.0.2: icmp_seq=9 ttl=62 time=0.891 ms
64 bytes from 23.0.0.2: icmp_seq=10 ttl=62 time=0.892 ms
64 bytes from 23.0.0.2: icmp_seq=11 ttl=62 time=0.893 ms
^C
--- bt.lboro ping statistics ---
11 packets transmitted, 11 received, 0% packet loss, time 10000ms
rtt min/avg/max/mdev = 0.745/0.861/0.893/0.045 ms
bt002@BT002:~$
```

(b) Laptop 2 (BT002)

```
bt003@BT003:~$ ping bt.lboro
PING bt.lboro (23.0.0.2) 56(84) bytes of data:
64 bytes from 23.0.0.2: icmp_seq=1 ttl=62 time=0.845 ms
64 bytes from 23.0.0.2: icmp_seq=2 ttl=62 time=0.884 ms
64 bytes from 23.0.0.2: icmp_seq=3 ttl=62 time=0.862 ms
64 bytes from 23.0.0.2: icmp_seq=4 ttl=62 time=0.868 ms
64 bytes from 23.0.0.2: icmp_seq=5 ttl=62 time=0.878 ms
64 bytes from 23.0.0.2: icmp_seq=6 ttl=62 time=0.763 ms
64 bytes from 23.0.0.2: icmp_seq=7 ttl=62 time=0.874 ms
64 bytes from 23.0.0.2: icmp_seq=8 ttl=62 time=0.882 ms
64 bytes from 23.0.0.2: icmp_seq=9 ttl=62 time=0.879 ms
64 bytes from 23.0.0.2: icmp_seq=10 ttl=62 time=0.862 ms
64 bytes from 23.0.0.2: icmp_seq=11 ttl=62 time=0.880 ms
^C
--- bt.lboro ping statistics ---
11 packets transmitted, 11 received, 0% packet loss, time 10010ms
rtt min/avg/max/mdev = 0.763/0.861/0.884/0.045 ms
bt003@BT003:~$
```

(c) Laptop 3 (BT003)

Figure 4.6: Connectivity to bt.lboro using ping on all 3 Laptops.

Under the unfortunate condition that the primary server (Laptop 3) is disconnected from the Internet, availability of DNS service in BT is also tested. As shown in Figure 4.7, DNS records are successfully retrieved from the secondary DNS server (Laptop 2, IPv4 Address: 23.0.0.2) despite the fact that the primary server is down.




Figure 4.7: Querying DNS Records for `bt.lboro` on Laptop 2 When Primary DNS Server (Laptop 3) Is Down.

4.2.4 Commentary

4.2.4.1 Problem: Querying IPv6 DNS Records

Using `dig` directly does not return IPv6 addresses for domains. It turns out that `dig` queries A records by default. To query IPv6 DNS records, `dig AAAA domain-name` should be used.

4.2.4.2 Problem: Reverse DNS Not Working

Using `dig -x 23.0.0.2` does not return any answer initially. It's later realized that the reverse DNS zone name has been wrongly set to `0.0.23.in-addr.arpa`.

To set up reverse DNS properly, the DNS zone should be `3.in-addr.arpa`. In addition, the first column for each DNS record inside file `db.23` should be the reverse host name (eg. `2.0.0` for IP address `23.0.0.2`).

4.3 World Wide Web Service

4.3.1 Design

A World Wide Web (WWW) service[9][10] in the network allows any terminal devices to access the deployed webpages. The service in BT Network is established after routers and three laptops have been configured and DNS service been set up.

For this lab, **apache2** package is chosen as the tool to establish the web server on Laptop 1 (BT001). Using the default settings of this package is enough and it has a specific folder static web pages are stored.

4.3.2 Implementation

Install **apache2** package on Laptop 1 (BT001, IPv4 Address: 23.0.0.2) using following the command.

```
1 sudo apt-get install apache2
```

A HTML file named **index.html** is created as a test webpage, as detailed in Figure 4.8.

```
1 <html>
2 <header><title>BT Network</title></header>
3 <body>
4 <h1>Welcome to BT Network |</h1>
5 A Trustworthy Internet Service Provider at Loughborough University.
6 </body>
7 </html>
```

Figure 4.8: Contents of HTML File Named **index.html**

Then, the HTML file is copied to the folder **/var/www/html**. This folder is used to deploy webpages on the server. And the meaning of **-r** is to cover the same name file.

```
1 cp -r index.html /var/www/html
```

Then, install **links** package which is a command-line Web browser in Linux by using following command.

```
1 sudo apt-get install links
```

Then, **links** is started and the web page is accessed through URL with the local IP address.

```
1 links http://23.0.0.2/
```

Finally, add the following lines to the forward DNS file `/etc/bind/db.bt.lboro` on Primary DNS Server at Laptop 3 (BT003). This enables the clients to browse the webpages by domain name `bt.lboro`.

```
1 @ IN A 23.0.0.2
2 @ IN AAAA 2001:2300:0:0::2
```

4.3.3 Evaluation

The web service is tested using the following command on Laptop 2, which is neither a DNS Server not a Web Server.

```
1 links http://bt.lboro/
```

As shown in Figure 4.9, the service has been successfully set up and is accessible in the network.




Figure 4.9: Web Service Provided at `http://bt.lboro` .

4.4 Email Service

4.4.1 Design

Email service[11] allows users to communicate with each other based on texts. In BT Network, it is set up using `exim4` package on Laptop 3 (BT003).

4.4.2 Implementation

On Laptop 3 (BT003), the package `exim4` is installed and configured. Important steps in the configuration are shown in 4.10.

```
1 apt-get install exim4
2 dpkg-reconfigure exim4-config
```

The full configuration of Email service is detailed in Appendix E.

4.4.3 Evaluation

For evaluation, the following command are used to send a mail to DT Network.


```
1 echo "mua." | sendmail -v mail@dt3.lboro
```

The mail can be seen arriving the destination in Figure 4.11.


4.4.4 Commentary

4.4.4.1 Problem: Setting Up Domain Name

At the very beginning, when we were configuring `exim4`, we just skipped the step and forgot to set up the domain name. Later, we found that the mail cannot be sent to other groups. The domain name is reset as `bt.lboro` to solve the problem.




(a) Setting Up Mail Name.



(b) Setting Up Domain Name.

Figure 4.10: Important Configuration Steps for `exim4`.



(a) SMTP Success Message Is Returned from DT Network on Laptop 3 (BT003).



(b) Email Is Received on DT Network's Side.

Figure 4.11: Email Can Be Seen Arrived on Both Sides.

Chapter 5

Discussion

5.1 Conclusions

Several conclusions can be drawn from this lab.

1. BT Network, a small Tier-2 ISP, has been built and well tested.
2. BT Network provides both intra-domain and inter-domain Internet connection to its users. It serves common Internet applications including Web, DNS and Email as well.
3. BT Network forms and implements business relationships with neighbouring ISPs.
4. Both IS-IS and BGP routing protocols can provide alternative route(s) to the destination when one of the physical links is down.

5.2 Further Work

For the future, the following improvements are being considered.

1. Implement the alternative **next-hop solution** instead of "passive interface" as in Section 3.1.5.2.
2. Fully test the implementation of BGP routing in IPv6. We are unable to test it as no neighbouring ISP has set up IPv6 BGP routing as far as we know.
3. Provide other Internet services such as Dynamic Host Configuration Protocol (DHCP)[12] and File Transfer Protocol (FTP)[13].

Chapter 6

Contributions

Our team is composed of 5 members from Loughborough University.

6.1 Group Leader: Zhihao DAI

In this lab, Zhihao DAI contributes to designing the architecture of the network, assigning IP addresses to interfaces, setting up BGP routing protocol as well as securing access to routers. Additionally, DAI is responsible for organising and formatting the whole report.

6.2 Technical Director: Yunsong ZHANG

Yunsong ZHANG contributes to designing the architecture of the network, setting up IS-IS routing protocol as well as DNS services.

6.3 Network Engineer: Huijing LEI

Huijing LEI contributes to setting up Web services. Additionally, LEI is responsible for References section in the report.

6.4 Network Engineer: Changrong CHEN

Changrong CHEN contributes to assigning IP addresses to interfaces and setting up Email services.

6.5 Network Engineer: Yan HUANG

Yan HUANG contributes to setting up Email services. Additionally, HUANG is responsible for Discussion section in the report.

REFERENCES

References

- [1] Internet Protocol. RFC 791, September 1981.
- [2] Bob Hinden and Dr. Steve E. Deering. Internet Protocol, Version 6 (IPv6) Specification. RFC 2460, December 1998.
- [3] Gary S. Malkin. RIP Version 2. RFC 2453, November 1998.
- [4] OSI IS-IS Intra-domain Routing Protocol. RFC 1142, February 1990.
- [5] John Moy. OSPF Version 2. RFC 2328, April 1998.
- [6] Yakov Rekhter, Susan Hares, and Tony Li. A Border Gateway Protocol 4 (BGP-4). RFC 4271, January 2006.
- [7] Chris M. Lonvick and Tatu Ylonen. The Secure Shell (SSH) Transport Layer Protocol. RFC 4253, January 2006.
- [8] Domain names - implementation and specification. RFC 1035, November 1987.
- [9] Tim Berners-Lee. Universal Resource Identifiers in WWW: A Unifying Syntax for the Expression of Names and Addresses of Objects on the Network as used in the World-Wide Web. RFC 1630, June 1994.
- [10] Henrik Frystyk Nielsen, Jeffrey Mogul, Larry M Masinter, Roy T. Fielding, Jim Gettys, Paul J. Leach, and Tim Berners-Lee. Hypertext Transfer Protocol – HTTP/1.1. RFC 2616, June 1999.
- [11] Pete Resnick. Internet Message Format. RFC 5322, October 2008.
- [12] Ralph Droms. Dynamic Host Configuration Protocol. RFC 2131, March 1997.
- [13] File Transfer Protocol. RFC 959, October 1985.

Appendix A

Login Details

A.1 Laptops

```
1 Laptop 1 (BT001):  
2   Username: bt001  
3   Password: Bt9049.4581  
4  
5 Laptop 2 (BT002):  
6   Username: bt002  
7   Password: Bt8717.0801  
8  
9 Laptop 3 (BT003):  
10  Username: bt003  
11  Password: Bt6941.6657
```

A.2 Routers

Routers can be remotely accessed without entering the password using public-key authentication (root required) on Laptop 1 (BT001). The password for each router is one of $a^2+2ab+b^2$, a^2-ab+b^2 or a^2+b^2 .

```
1 Router 1 (BT-R001):  
2   Username: r001  
3  
4 Router 2 (BT-R002):  
5   Username: r002  
6  
7 Router 3 (BT-R003):  
8   Username: r003
```

Appendix B

Routers Configuration

B.1 Router 1 Configuration

Listing B.1: Contents of Configuration on Router 1 (BT-R001).

```
1
2 !
3 ! Last configuration change at 14:43:27 UTC Mon Nov 4 2019 by r001
4 !
5 version 15.0
6 service timestamps debug datetime msec
7 service timestamps log datetime msec
8 no service password-encryption
9 !
10 hostname BT-R001
11 !
12 boot-start-marker
13 boot-end-marker
14 !
15 !
16 no aaa new-model
17 dot11 syslog
18 ip source-route
19 !
20 !
21 !
22 !
23 ip cef
24 no ip domain lookup
25 ip domain name bt.lboro
26 ipv6 unicast-routing
```


APPENDIX B. ROUTERS CONFIGURATION

```
27 ipv6 cef
28 multilink bundle-name authenticated
29 !
30 !
31 !
32 !
33 !
34 !
35 !
36 !
37 !
38 voice-card 0
39 !
40 !
41 !
42 !
43 !
44 license udi pid CISC02801 sn FCZ1339C10B
45 username r001 privilege 15 secret 5 $1$zzE2$UaRCPrx0iQCcbQ1Jg1Ys21
46 !
47 !
48 ip ssh version 2
49 ip ssh dh min size 4096
50 ip ssh pubkey-chain
51   username r001
52     key-hash ssh-rsa AFC27BDF03A4FB6173D4D0482B4C084A
53     key-hash ssh-rsa 601E98BF47BA6CC0089AEC1177F6DEE5 root@BT001
54 !
55 !
56 !
57 !
58 !
59 interface Loopback0
60   ip address 23.0.255.1 255.255.255.255
61   ipv6 address 2001:2300:FFFF:1::/128
62 !
63 interface FastEthernet0/0
64   ip address 23.0.0.49 255.255.255.252
65   ip router isis
66   duplex auto
67   speed auto
68   ipv6 address 2001:2300:0:3::1/64
69   ipv6 router isis
70 !
71 interface FastEthernet0/1
72   ip address 23.0.0.57 255.255.255.252
73   ip router isis
```

APPENDIX B. ROUTERS CONFIGURATION

```
74 duplex auto
75 speed auto
76 ipv6 address 2001:2300:0:5::1/64
77 ipv6 router isis
78 !
79 interface FastEthernet0/1/0
80 ip address 23.0.0.1 255.255.255.240
81 ip router isis
82 duplex auto
83 speed auto
84 ipv6 address 2001:2300::1/64
85 ipv6 router isis
86 !
87 router isis
88 net 49.0001.0230.0025.5001.00
89 is-type level-1
90 passive-interface Loopback0
91 !
92 router bgp 2030
93 no synchronization
94 bgp log-neighbor-changes
95 network 23.0.0.0 mask 255.255.255.240
96 aggregate-address 23.0.0.0 255.0.0.0 summary-only
97 neighbor 23.0.255.2 remote-as 2030
98 neighbor 23.0.255.2 update-source Loopback0
99 neighbor 23.0.255.3 remote-as 2030
100 neighbor 23.0.255.3 update-source Loopback0
101 neighbor 2001:2300:FFFF:2:: remote-as 2030
102 neighbor 2001:2300:FFFF:2:: update-source Loopback0
103 neighbor 2001:2300:FFFF:3:: remote-as 2030
104 neighbor 2001:2300:FFFF:3:: update-source Loopback0
105 no auto-summary
106 !
107 address-family ipv6
108 network 2001:2300::/64
109 aggregate-address 2001:2300::/32 summary-only
110 neighbor 2001:2300:FFFF:2:: activate
111 neighbor 2001:2300:FFFF:3:: activate
112 exit-address-family
113 !
114 ip forward-protocol nd
115 !
116 !
117 no ip http server
118 no ip http secure-server
119 !
120 !
```

APPENDIX B. ROUTERS CONFIGURATION

```
121 !
122 !
123 control-plane
124 !
125 !
126 !
127 mgcp fax t38 ecm
128 mgcp behavior g729-variants static-pt
129 !
130 !
131 !
132 !
133 line con 0
134 line aux 0
135 line vty 0 4
136   login local
137   transport input telnet ssh
138 !
139 scheduler allocate 20000 1000
140 end
```

B.2 Router 2 Configuration

Listing B.2: Contents of Configuration on Router 2 (BT-R002).

```
1
2 !
3 ! Last configuration change at 14:35:53 UTC Mon Nov 4 2019 by r002
4 !
5 version 15.0
6 service timestamps debug datetime msec
7 service timestamps log datetime msec
8 no service password-encryption
9 !
10 hostname BT-R002
11 !
12 boot-start-marker
13 boot-end-marker
14 !
15 !
16 no aaa new-model
17 dot11 syslog
18 ip source-route
19 !
20 !
21 !
22 !
23 ip cef
24 no ip domain lookup
25 ip domain name bt.lboro
26 ipv6 unicast-routing
27 ipv6 cef
28 multilink bundle-name authenticated
29 !
30 !
31 !
32 !
33 !
34 !
35 !
36 !
37 !
38 voice-card 0
39 !
40 !
41 !
42 !
```

APPENDIX B. ROUTERS CONFIGURATION

```
43 !
44 license udi pid CISC02801 sn FCZ1339C100
45 username r002 privilege 15 secret 5 $1$..EF$yfYESFK61lTyMnP.f1ABh0
46 !
47 !
48 ip ssh version 2
49 ip ssh dh min size 4096
50 ip ssh pubkey-chain
51   username r002
52     key-hash ssh-rsa 601E98BF47BA6CC0089AEC1177F6DEE5 root@BT001
53 !
54 !
55 !
56 !
57 !
58 interface Loopback0
59   ip address 23.0.255.2 255.255.255.255
60   ipv6 address 2001:2300:FFFF:2::/128
61 !
62 interface FastEthernet0/0
63   ip address 23.0.0.50 255.255.255.252
64   ip router isis
65   duplex auto
66   speed auto
67   ipv6 address 2001:2300:0:3::2/64
68   ipv6 router isis
69 !
70 interface FastEthernet0/1
71   ip address 23.0.0.53 255.255.255.252
72   ip router isis
73   duplex auto
74   speed auto
75   ipv6 address 2001:2300:0:4::1/64
76   ipv6 router isis
77 !
78 interface FastEthernet0/1/0
79 !
80 interface FastEthernet0/1/1
81   switchport access vlan 3
82 !
83 interface FastEthernet0/1/2
84 !
85 interface FastEthernet0/1/3
86 !
87 interface Vlan1
88   ip address 23.0.0.17 255.255.255.240
89   ip router isis
```

APPENDIX B. ROUTERS CONFIGURATION

```
90  ipv6 address 2001:2300:0:1::1/64
91  ipv6 router isis
92  !
93  interface Vlan3
94  ip address 23.0.0.61 255.255.255.252
95  ipv6 address 2001:2300:0:6::1/64
96  !
97  router isis
98  net 49.0001.0230.0025.5002.00
99  is-type level-1
100 passive-interface Vlan3
101 passive-interface Loopback0
102 !
103 router bgp 2030
104  bgp log-neighbor-changes
105  neighbor 23.0.0.62 remote-as 3040
106  neighbor 23.0.255.1 remote-as 2030
107  neighbor 23.0.255.1 update-source Loopback0
108  neighbor 23.0.255.3 remote-as 2030
109  neighbor 23.0.255.3 update-source Loopback0
110  neighbor 2001:2300:0:6::2 remote-as 3040
111  neighbor 2001:2300:FFFF:1:: remote-as 2030
112  neighbor 2001:2300:FFFF:1:: update-source Loopback0
113  neighbor 2001:2300:FFFF:3:: remote-as 2030
114  neighbor 2001:2300:FFFF:3:: update-source Loopback0
115  !
116  address-family ipv4
117    no synchronization
118    network 23.0.0.16 mask 255.255.255.240
119    aggregate-address 23.0.0.0 255.0.0.0 summary-only
120    neighbor 23.0.0.62 activate
121    neighbor 23.0.255.1 activate
122    neighbor 23.0.255.3 activate
123    neighbor 2001:2300:0:6::2 activate
124    neighbor 2001:2300:FFFF:1:: activate
125    neighbor 2001:2300:FFFF:3:: activate
126    no auto-summary
127  exit-address-family
128  !
129  address-family ipv6
130    network 2001:2300:0:1::/64
131    aggregate-address 2001:2300::/32 summary-only
132    neighbor 2001:2300:0:6::2 activate
133    neighbor 2001:2300:FFFF:1:: activate
134    neighbor 2001:2300:FFFF:3:: activate
135  exit-address-family
136  !
```

APPENDIX B. ROUTERS CONFIGURATION

```
137 ip forward-protocol nd
138 !
139 !
140 no ip http server
141 no ip http secure-server
142 !
143 !
144 !
145 !
146 control-plane
147 !
148 !
149 !
150 mgcp fax t38 ecm
151 mgcp behavior g729-variants static-pt
152 !
153 !
154 !
155 !
156 line con 0
157 line aux 0
158 line vty 0 4
159   login local
160   transport input telnet ssh
161 !
162 scheduler allocate 20000 1000
163 end
```

B.3 Router 3 Configuration

Listing B.3: Contents of Configuration on Router 3 (BT-R003).

```
1
2 !
3 ! Last configuration change at 06:40:58 UTC Sun Nov 3 2019 by r003
4 !
5 version 15.0
6 service timestamps debug datetime msec
7 service timestamps log datetime msec
8 no service password-encryption
9 !
10 hostname BT-R003
11 !
12 boot-start-marker
13 boot-end-marker
14 !
15 !
16 no aaa new-model
17 dot11 syslog
18 ip source-route
19 !
20 !
21 !
22 !
23 ip cef
24 no ip domain lookup
25 ip domain name bt.lboro
26 ipv6 unicast-routing
27 ipv6 cef
28 multilink bundle-name authenticated
29 !
30 !
31 !
32 !
33 !
34 !
35 !
36 !
37 !
38 voice-card 0
39 !
40 !
41 !
42 !
```


APPENDIX B. ROUTERS CONFIGURATION

```
43 !
44 license udi pid CISC02801 sn FCZ124112JK
45 username r003 privilege 15 secret 5 $1$Jn3f$1SSIZA5X0c0JxuXDkz8rE/
46 !
47 !
48 ip ssh version 2
49 ip ssh dh min size 4096
50 ip ssh pubkey-chain
51   username r003
52     key-hash ssh-rsa 601E98BF47BA6CC0089AEC1177F6DEE5 root@BT001
53 !
54 !
55 !
56 !
57 !
58 interface Loopback0
59   ip address 23.0.255.3 255.255.255.255
60   ipv6 address 2001:2300:FFFF:3::/128
61 !
62 interface FastEthernet0/0
63   ip address 23.0.0.58 255.255.255.252
64   ip router isis
65   duplex auto
66   speed auto
67   ipv6 address 2001:2300:0:5::2/64
68   ipv6 router isis
69 !
70 interface FastEthernet0/1
71   ip address 23.0.0.54 255.255.255.252
72   ip router isis
73   duplex auto
74   speed auto
75   ipv6 address 2001:2300:0:4::2/64
76   ipv6 router isis
77 !
78 interface FastEthernet0/1/0
79   switchport access vlan 2
80 !
81 interface FastEthernet0/1/1
82   switchport access vlan 5
83 !
84 interface FastEthernet0/1/2
85   switchport access vlan 4
86 !
87 interface FastEthernet0/1/3
88 !
89 interface Vlan1
```

APPENDIX B. ROUTERS CONFIGURATION

```
90 no ip address
91 !
92 interface Vlan2
93 ip address 23.0.0.33 255.255.255.240
94 ip router isis
95 ipv6 address 2001:2300:0:2::1/64
96 ipv6 router isis
97 !
98 interface Vlan4
99 ip address 56.0.0.62 255.255.255.252
100 ipv6 address 2001:5600:0:6::2/64
101 !
102 interface Vlan5
103 ip address 100.100.2.2 255.255.255.252
104 !
105 router isis
106 net 49.0001.0230.0025.5003.00
107 is-type level-1
108 passive-interface Vlan4
109 passive-interface Vlan5
110 passive-interface Loopback0
111 !
112 router bgp 2030
113 no synchronization
114 bgp log-neighbor-changes
115 network 23.0.0.32 mask 255.255.255.240
116 aggregate-address 23.0.0.0 255.0.0.0 summary-only
117 neighbor 23.0.255.1 remote-as 2030
118 neighbor 23.0.255.1 update-source Loopback0
119 neighbor 23.0.255.2 remote-as 2030
120 neighbor 23.0.255.2 update-source Loopback0
121 neighbor 2001:2300:FFFF:1:: remote-as 2030
122 neighbor 2001:2300:FFFF:1:: update-source Loopback0
123 neighbor 2001:2300:FFFF:2:: remote-as 2030
124 neighbor 2001:2300:FFFF:2:: update-source Loopback0
125 neighbor 2001:5600:0:6::1 remote-as 5060
126 neighbor 56.0.0.61 remote-as 5060
127 neighbor 56.0.0.61 filter-list 1 out
128 neighbor 100.100.2.1 remote-as 42
129 neighbor 100.100.2.1 filter-list 1 out
130 no auto-summary
131 !
132 address-family ipv6
133 network 2001:2300:0:2::/64
134 aggregate-address 2001:2300::/32 summary-only
135 neighbor 2001:2300:FFFF:1:: activate
136 neighbor 2001:2300:FFFF:2:: activate
```

APPENDIX B. ROUTERS CONFIGURATION

```
137  exit-address-family
138  !
139  ip forward-protocol nd
140  !
141  ip as-path access-list 1 deny _42_
142  ip as-path access-list 1 deny _5060_
143  ip as-path access-list 1 permit .*
144  !
145  no ip http server
146  no ip http secure-server
147  !
148  !
149  !
150  !
151  control-plane
152  !
153  !
154  !
155  mgcp fax t38 ecm
156  mgcp behavior g729-variants static-pt
157  !
158  !
159  !
160  !
161  line con 0
162  line aux 0
163  line vty 0 4
164  login local
165  transport input telnet ssh
166  !
167  scheduler allocate 20000 1000
168  end
```

Appendix C

Laptops Configuration

C.1 Laptop 1 Configuration

Listing C.1: Contents of Interfaces Network Configuration File Located at `/etc/network/interfaces` on Laptop 1 (BT001).

```
1 # This file describes the network interfaces available on your system
2 # and how to activate them. For more information, see interfaces(5).
3
4 # The loopback network interface
5 auto lo
6 iface lo inet loopback
7
8 # The primary network interface
9 auto eth0
10 iface eth0 inet static
11     address 23.0.0.2
12     netmask 255.255.255.240
13     gateway 23.0.0.1
14     dns-nameservers 23.0.0.34 23.0.0.2
15 # up flush-mail
16
17
18 iface eth0 inet6 static
19     address 2001:2300::2
20     netmask 64
21     gateway 2001:2300::1
```

C.2 Laptop 2 Configuration

Listing C.2: Contents of Interfaces Network Configuration File Located at `/etc/network/interfaces` on Laptop 2 (BT002).

```
1 # This file describes the network interfaces available on your system
2 # and how to activate them. For more information, see interfaces(5).
3
4 # The loopback network interface
5 auto lo
6 iface lo inet loopback
7
8 # The primary network interface
9 auto eth0
10 iface eth0 inet static
11     address 23.0.0.18
12     netmask 255.255.255.240
13     gateway 23.0.0.17
14     dns-nameservers 23.0.0.34 23.0.0.2
15 #     up flush-mail
16
17 iface eth0 inet6 static
18     address 2001:2300:0:1::2
19     netmask 64
20     gateway 2001:2300:0:1::1
```

C.3 Laptop 3 Configuration

Listing C.3: Contents of Interfaces Network Configuration File Located at `/etc/network/interfaces` on Laptop 3 (BT003).

```
1 # This file describes the network interfaces available on your system
2 # and how to activate them. For more information, see interfaces(5).
3
4 # The loopback network interface
5 auto lo
6 iface lo inet loopback
7
8 # The primary network interface
9 auto eth0
10 iface eth0 inet static
11     address 23.0.0.2
12     netmask 255.255.255.240
13     gateway 23.0.0.1
14     dns-nameservers 23.0.0.34 23.0.0.2
15 # up flush-mail
16
17
18 iface eth0 inet6 static
19     address 2001:2300::2
20     netmask 64
21     gateway 2001:2300::1
```

Appendix D

DNS Configuration

D.1 Primary DNS Server (Laptop 3)

Listing D.1: Contents of DNS Configuration File Located at `/etc/bind/named.conf.options` on Primary DNS Server (Laptop 3).

```
1 options {
2     directory "/var/cache/bind";
3
4     // If there is a firewall between you and nameservers you want
5     // to talk to, you may need to fix the firewall to allow multiple
6     // ports to talk. See http://www.kb.cert.org/vuls/id/800113
7
8     // If your ISP provided one or more IP addresses for stable
9     // nameservers, you probably want to use them as forwarders.
10    // Uncomment the following block, and insert the addresses replacing
11    // the all-0's placeholder.
12
13    // forwarders {
14    //     0.0.0.0;
15    // };
16
17    //=====
18    // If BIND logs error messages about the root key being expired,
19    // you will need to update your keys. See https://www.isc.org/bind-keys
20    //=====
21    // dnssec-validation auto;
22
23    auth-nxdomain no;          # conform to RFC1035
24    listen-on-v6 { any; };
```

APPENDIX D. DNS CONFIGURATION

```
25 forwarders { 10.2.2.1; };
26 recursion yes;
27 dnssec-enable no;
28 dnssec-validation no;
29 empty-zones-enable no;
30 allow-query {any;};
31 };
```

Listing D.2: Contents of Local DNS Configuration File Located at `/etc/bind/named.conf.local` on Primary DNS Server (Laptop 3).

```
1 //
2 // Do any local configuration here
3 //
4
5 // Consider adding the 1918 zones here, if they are not used in your
6 // organization
7 //include "/etc/bind/zones.rfc1918";
8
9 zone "bt.lboro" {
10     type master;
11     file "/etc/bind/db.bt.lboro";
12     allow-transfer { 23.0.0.2; };
13     also-notify { 23.0.0.2; };
14 };
15
16 zone "23.in-addr.arpa" {
17     type master;
18     file "/etc/bind/db.23";
19     allow-transfer { 23.0.0.2; };
20     also-notify { 23.0.0.2; };
21 };
```


D.2 Secondary DNS Server (Laptop 1)

Listing D.3: Contents of DNS Configuration File Located at `/etc/bind/named.conf.options` on Secondary DNS Server (Laptop 1).

```
1 options {
2     directory "/var/cache/bind";
3
4     // If there is a firewall between you and nameservers you want
5     // to talk to, you may need to fix the firewall to allow multiple
6     // ports to talk. See http://www.kb.cert.org/vuls/id/800113
7
8     // If your ISP provided one or more IP addresses for stable
9     // nameservers, you probably want to use them as forwarders.
10    // Uncomment the following block, and insert the addresses replacing
11    // the all-0's placeholder.
12
13    // forwarders {
14    //     0.0.0.0;
15    // };
16
17    //=====
18    // If BIND logs error messages about the root key being expired,
19    // you will need to update your keys. See https://www.isc.org/bind-keys
20    //=====
21    dnssec-validation auto;
22
23    auth-nxdomain no;      # conform to RFC1035
24    listen-on-v6 { any; };
25};
```

Listing D.4: Contents of Local DNS Configuration File Located at `/etc/bind/named.conf.local` on Secondary DNS Server (Laptop 1).

```
1 //
2 // Do any local configuration here
3 //
4
5 // Consider adding the 1918 zones here, if they are not used in your
6 // organization
7 //include "/etc/bind/zones.rfc1918";
8
9 zone "bt.lboro" {
10     type slave;
11     file "db.bt.lboro";
12     masters { 23.0.0.34; };
```

APPENDIX D. DNS CONFIGURATION

```
13 };  
14  
15 zone "0.0.23.in-addr.arpa" {  
16     type slave;  
17     file "db.23";  
18     masters { 23.0.0.34; };  
19 };
```

Appendix E

Email Configuration

Listing E.1: Contents of Generated Email Configuration File Located at `/var/lib/exim4/config.autogenerated` on Email Server (Laptop 3).

```
1 #####
2 # WARNING WARNING WARNING
3 # WARNING WARNING WARNING
4 # WARNING WARNING WARNING
5 # WARNING WARNING WARNING
6 # WARNING WARNING WARNING
7 # This file was generated dynamically from
8 # non-split config (/etc/exim4/exim4.conf.localmacros
9 # and /etc/exim4/exim4.conf.template).
10 # The config files are supplemented with package installation/configuration
11 # settings managed by debconf. This data is stored in
12 # /etc/exim4/update-exim4.conf.conf
13 # Any changes you make here will be lost.
14 # See /usr/share/doc/exim4-base/README.Debian.gz and update-exim4.conf(8)
15 # for instructions of customization.
16 # WARNING WARNING WARNING
17 # WARNING WARNING WARNING
18 # WARNING WARNING WARNING
19 # WARNING WARNING WARNING
20 # WARNING WARNING WARNING
21 #####
22
23 exim_path = /usr/sbin/exim4
24
25 .ifndef CONFDIR
26 CONFDIR = /etc/exim4
27 .endif
```

APPENDIX E. EMAIL CONFIGURATION

```
28
29 UPEX4CmacrosUPEX4C = 1
30 #####
31 # the following macro definitions were created
32 # dynamically by /usr/sbin/update-exim4.conf
33 .ifndef MAIN_PACKAGE_VERSION
34 MAIN_PACKAGE_VERSION=4.82-3ubuntu2
35 .endif
36 .ifndef MAIN_LOCAL_DOMAINS
37 MAIN_LOCAL_DOMAINS=@:localhost:bt003:bt.lboro
38 .endif
39 .ifndef MAIN_RELAY_TO_DOMAINS
40 MAIN_RELAY_TO_DOMAINS=empty
41 .endif
42 .ifndef ETC_MAILNAME
43 ETC_MAILNAME=bt.lboro
44 .endif
45 .ifndef LOCAL_DELIVERY
46 LOCAL_DELIVERY=mail_spool
47 .endif
48 .ifndef MAIN_RELAY_NETS
49 MAIN_RELAY_NETS=: 127.0.0.1 : :::1
50 .endif
51 .ifndef DCreadhost
52 DCreadhost=empty
53 .endif
54 .ifndef DCsmarthost
55 DCsmarthost=empty
56 .endif
57 .ifndef DC_eximconfig_configtype
58 DC_eximconfig_configtype=internet
59 .endif
60 .ifndef DCconfig_internet
61 DCconfig_internet=1
62 .endif
63 #####
64
65
66 domainlist local_domains = MAIN_LOCAL_DOMAINS
67
68 domainlist relay_to_domains = MAIN_RELAY_TO_DOMAINS
69
70 hostlist relay_from_hosts = MAIN_RELAY_NETS
71
72 .ifndef MAIN_PRIMARY_HOSTNAME_AS_QUALIFY_DOMAIN
73 .ifndef MAIN_QUALIFY_DOMAIN
74 qualify_domain = ETC_MAILNAME
```

APPENDIX E. EMAIL CONFIGURATION

```
75 .else
76 qualify_domain = MAIN_QUALIFY_DOMAIN
77 .endif
78 .endif
79
80 .ifdef MAIN_LOCAL_INTERFACES
81 local_interfaces = MAIN_LOCAL_INTERFACES
82 .endif
83
84 .ifndef LOCAL_DELIVERY
85 LOCAL_DELIVERY=mail_spool
86 .endif
87
88 gecos_pattern = ^([^,:]*)
89 gecos_name = $1
90
91 .ifndef CHECK_RCPT_LOCAL_LOCALPARTS
92 CHECK_RCPT_LOCAL_LOCALPARTS = ^[.] : ^.*[!@%/'#&?]
93 .endif
94
95 .ifndef CHECK_RCPT_REMOTE_LOCALPARTS
96 CHECK_RCPT_REMOTE_LOCALPARTS = ^[./|] : ^.*[!@%/'#&?] : ^.*[\\.\|].*/
97 .endif
98
99 .ifndef MAIN_LOG_SELECTOR
100 MAIN_LOG_SELECTOR = +tls_peerdn
101 .endif
102
103 .ifndef MAIN_ACL_CHECK_MAIL
104 MAIN_ACL_CHECK_MAIL = acl_check_mail
105 .endif
106 acl_smtp_mail = MAIN_ACL_CHECK_MAIL
107
108 .ifndef MAIN_ACL_CHECK_RCPT
109 MAIN_ACL_CHECK_RCPT = acl_check_rcpt
110 .endif
111 acl_smtp_rcpt = MAIN_ACL_CHECK_RCPT
112
113 .ifndef MAIN_ACL_CHECK_DATA
114 MAIN_ACL_CHECK_DATA = acl_check_data
115 .endif
116 acl_smtp_data = MAIN_ACL_CHECK_DATA
117
118 .ifdef MESSAGE_SIZE_LIMIT
119 message_size_limit = MESSAGE_SIZE_LIMIT
120 .endif
121
```

APPENDIX E. EMAIL CONFIGURATION

```
122 .ifdef MAIN_ALLOW_DOMAIN_LITERALS
123 allow_domain_literals
124 .endif
125
126 .ifndef DC_minimaldns
127 .ifndef MAIN_HOST_LOOKUP
128 MAIN_HOST_LOOKUP = *
129 .endif
130 host_lookup = MAIN_HOST_LOOKUP
131 .endif
132
133 .ifdef MAIN_HARDCODE_PRIMARY_HOSTNAME
134 primary_hostname = MAIN_HARDCODE_PRIMARY_HOSTNAME
135 .endif
136
137 .ifdef MAIN_SMTP_ACCEPT_MAX_NOMAIL_HOSTS
138 smtp_accept_max_nonmail_hosts = MAIN_SMTP_ACCEPT_MAX_NOMAIL_HOSTS
139 .endif
140
141 .ifndef MAIN_FORCE_SENDER
142 local_from_check = false
143 local_sender_retain = true
144 untrusted_set_sender = *
145 .endif
146
147 .ifndef MAIN_IGNORE_BOUNCE_ERRORS_AFTER
148 MAIN_IGNORE_BOUNCE_ERRORS_AFTER = 2d
149 .endif
150 ignore_bounce_errors_after = MAIN_IGNORE_BOUNCE_ERRORS_AFTER
151
152 .ifndef MAIN_TIMEOUT_FROZEN_AFTER
153 MAIN_TIMEOUT_FROZEN_AFTER = 7d
154 .endif
155 timeout_frozen_after = MAIN_TIMEOUT_FROZEN_AFTER
156
157 .ifndef MAIN_FREEZE_TELL
158 MAIN_FREEZE_TELL = postmaster
159 .endif
160 freeze_tell = MAIN_FREEZE_TELL
161
162 .ifndef SPOOLDIR
163 SPOOLDIR = /var/spool/exim4
164 .endif
165 spool_directory = SPOOLDIR
166
167 .ifndef MAIN_TRUSTED_USERS
168 MAIN_TRUSTED_USERS = uucp
```

APPENDIX E. EMAIL CONFIGURATION

```
169 .endif
170 trusted_users = MAIN_TRUSTED_USERS
171 .ifdef MAIN_TRUSTED_GROUPS
172 trusted_groups = MAIN_TRUSTED_GROUPS
173 .endif
174
175 .ifdef MAIN_TLS_ENABLE
176 .ifndef MAIN_TLS_ADVERTISE_HOSTS
177 MAIN_TLS_ADVERTISE_HOSTS = *
178 .endif
179 tls_advertise_hosts = MAIN_TLS_ADVERTISE_HOSTS
180
181 .ifdef MAIN_TLS_CERTKEY
182 tls_certificate = MAIN_TLS_CERTKEY
183 .else
184 .ifndef MAIN_TLS_CERTIFICATE
185 MAIN_TLS_CERTIFICATE = CONFDIR/exim.crt
186 .endif
187 tls_certificate = MAIN_TLS_CERTIFICATE
188
189 .ifndef MAIN_TLS_PRIVATEKEY
190 MAIN_TLS_PRIVATEKEY = CONFDIR/exim.key
191 .endif
192 tls_privatekey = MAIN_TLS_PRIVATEKEY
193 .endif
194
195 .ifndef MAIN_TLS_VERIFY_CERTIFICATES
196 MAIN_TLS_VERIFY_CERTIFICATES = ${if exists{/etc/ssl/certs/ca-certificates.crt
    }\
197                                     {/etc/ssl/certs/ca-certificates.crt}\
198                                     {/dev/null}}
199 .endif
200 tls_verify_certificates = MAIN_TLS_VERIFY_CERTIFICATES
201
202 .ifdef MAIN_TLS_VERIFY_HOSTS
203 tls_verify_hosts = MAIN_TLS_VERIFY_HOSTS
204 .endif
205
206 .ifdef MAIN_TLS_TRY_VERIFY_HOSTS
207 tls_try_verify_hosts = MAIN_TLS_TRY_VERIFY_HOSTS
208 .endif
209
210 .endif
211
212 .ifdef MAIN_LOG_SELECTOR
213 log_selector = MAIN_LOG_SELECTOR
214 .endif
```

APPENDIX E. EMAIL CONFIGURATION

```
215
216 begin acl
217
218 acl_local_deny_exceptions:
219     accept
220         hosts = ${if exists{CONFDIR/host_local_deny_exceptions}\
221                 {CONFDIR/host_local_deny_exceptions}\
222                 {}}
223     accept
224         senders = ${if exists{CONFDIR/sender_local_deny_exceptions}\
225                   {CONFDIR/sender_local_deny_exceptions}\
226                   {}}
227     accept
228         hosts = ${if exists{CONFDIR/local_host_whitelist}\
229                 {CONFDIR/local_host_whitelist}\
230                 {}}
231     accept
232         senders = ${if exists{CONFDIR/local_sender_whitelist}\
233                   {CONFDIR/local_sender_whitelist}\
234                   {}}
235
236 .ifdef LOCAL_DENY_EXCEPTIONS_LOCAL_ACL_FILE
237 .include LOCAL_DENY_EXCEPTIONS_LOCAL_ACL_FILE
238 .endif
239
240 .ifdef WHITELIST_LOCAL_DENY_LOCAL_ACL_FILE
241 .include WHITELIST_LOCAL_DENY_LOCAL_ACL_FILE
242 .endif
243
244 acl_check_mail:
245     .ifdef CHECK_MAIL_HELO_ISSUED
246     deny
247         message = no HELO given before MAIL command
248         condition = ${if def:sender_helo_name {no}{yes}}
249     .endif
250
251     accept
252
253 acl_check_rcpt:
254
255     accept
256         hosts = :
257         control = dkim_disable_verify
258
259 .ifdef DC_minimaldns
260     warn
261         control = dkim_disable_verify
```


APPENDIX E. EMAIL CONFIGURATION

```
262 .else
263 .ifdef DISABLE_DKIM_VERIFY
264     warn
265     control = dkim_disable_verify
266 .endif
267 .endif
268
269 .ifdef CHECK_RCPT_LOCAL_LOCALPARTS
270     deny
271     domains = +local_domains
272     local_parts = CHECK_RCPT_LOCAL_LOCALPARTS
273     message = restricted characters in address
274 .endif
275
276
277
278 .ifdef CHECK_RCPT_REMOTE_LOCALPARTS
279     deny
280     domains = !+local_domains
281     local_parts = CHECK_RCPT_REMOTE_LOCALPARTS
282     message = restricted characters in address
283 .endif
284
285     accept
286     .ifndef CHECK_RCPT_POSTMASTER
287         local_parts = postmaster
288     .else
289         local_parts = CHECK_RCPT_POSTMASTER
290     .endif
291     domains = +local_domains : +relay_to_domains
292
293 .ifdef CHECK_RCPT_VERIFY_SENDER
294     deny
295     message = Sender verification failed
296     !acl = acl_local_deny_exceptions
297     !verify = sender
298 .endif
299
300     deny
301     !acl = acl_local_deny_exceptions
302     senders = ${if exists{CONFFDIR/local_sender_callout}\
303                 {CONFFDIR/local_sender_callout}\
304                 {}}
305     !verify = sender/callout
306
307     accept
308     hosts = +relay_from_hosts
```

APPENDIX E. EMAIL CONFIGURATION

```
309     control = submission/sender_retain
310     control = dkim_disable_verify
311
312     accept
313         authenticated = *
314         control = submission/sender_retain
315         control = dkim_disable_verify
316
317     require
318         message = relay not permitted
319         domains = +local_domains : +relay_to_domains
320
321     require
322         verify = recipient
323
324     deny
325         !acl = acl_local_deny_exceptions
326         recipients = ${if exists{CONFFDIR/local_rcpt_callout}\
327                     {CONFFDIR/local_rcpt_callout}\
328                     {}}
329         !verify = recipient/callout
330
331     deny
332         message = sender envelope address $sender_address is locally blacklisted
333         here. If you think this is wrong, get in touch with postmaster
334         !acl = acl_local_deny_exceptions
335         senders = ${if exists{CONFFDIR/local_sender_blacklist}\
336                 {CONFFDIR/local_sender_blacklist}\
337                 {}}
338
339     deny
340         message = sender IP address $sender_host_address is locally blacklisted
341         here. If you think this is wrong, get in touch with postmaster
342         !acl = acl_local_deny_exceptions
343         hosts = ${if exists{CONFFDIR/local_host_blacklist}\
344                {CONFFDIR/local_host_blacklist}\
345                {}}
346
347     .ifdef CHECK_RCPT_REVERSE_DNS
348     warn
349         condition = ${if and{{def:sender_host_address}{!def:sender_host_name}}\
350                     {yes}{no}}
351         add_header = X-Host-Lookup-Failed: Reverse DNS lookup failed for
352                     $sender_host_address (${if eq{$host_lookup_failed}{1}{failed}{deferred}})
353     .endif
354
355     .ifdef CHECK_RCPT_SPF
```

APPENDIX E. EMAIL CONFIGURATION

```
353 deny
354     message = [SPF] $sender_host_address is not allowed to send mail from \
355               ${if def:sender_address_domain ${sender_address_domain}{
356               $sender_helo_name}}. \
357               Please see \
358               http://www.openspf.org/Why?scope=${if def:sender_address_domain \
359               {mfrom}{helo}};identity=${if def:sender_address_domain \
360               ${sender_address}${sender_helo_name}};ip=$sender_host_address
361     log_message = SPF check failed.
362     !acl = acl_local_deny_exceptions
363     condition = ${run{/usr/bin/spfquery.mail-spf-perl --ip \
364                 ${quote:$sender_host_address} --identity \
365                 ${if def:sender_address_domain \
366                 {--scope mfrom --identity ${quote:$sender_address}}\
367                 {--scope helo --identity ${quote:$sender_helo_name}}}}\
368                 {no}${if eq ${runrc}{1}{yes}{no}}}
369
370 defer
371     message = Temporary DNS error while checking SPF record. Try again later.
372     !acl = acl_local_deny_exceptions
373     condition = ${if eq ${runrc}{5}{yes}{no}}
374
375 warn
376     condition = ${if <=${runrc}{6}{yes}{no}}
377     add_header = Received-SPF: ${if eq ${runrc}{0}{pass}\
378                               ${if eq ${runrc}{2}{softfail}\
379                               ${if eq ${runrc}{3}{neutral}\
380                               ${if eq ${runrc}{4}{permerror}\
381                               ${if eq ${runrc}{6}{none}{error}}}}}}}\
382     } client-ip=$sender_host_address; \
383     ${if def:sender_address_domain \
384     {envelope-from=${sender_address}; }}}\
385     helo=$sender_helo_name
386
387 warn
388     log_message = Unexpected error in SPF check.
389     condition = ${if >${runrc}{6}{yes}{no}}
390 .endif
391
392 .ifdef CHECK_RCPT_IP_DNSBLS
393 warn
394     dnslists = CHECK_RCPT_IP_DNSBLS
395     add_header = X-Warning: $sender_host_address is listed at $dnslist_domain
396     ($dnslist_value: $dnslist_text)
397     log_message = $sender_host_address is listed at $dnslist_domain (
398     $dnslist_value: $dnslist_text)
399 .endif
```

APPENDIX E. EMAIL CONFIGURATION

```
397
398 .ifdef CHECK_RCPT_DOMAIN_DNSBLS
399 warn
400     !senders = ${if exists{CONFFDIR/local_domain_dnsbl_whitelist}\
401                 {CONFFDIR/local_domain_dnsbl_whitelist}\
402                 {}}
403     dnslists = CHECK_RCPT_DOMAIN_DNSBLS
404     add_header = X-Warning: $sender_address_domain is listed at
405                 $dnslist_domain ($dnslist_value: $dnslist_text)
406     log_message = $sender_address_domain is listed at $dnslist_domain (
407                 $dnslist_value: $dnslist_text)
408 .endif
409
410 .ifdef CHECK_RCPT_LOCAL_ACL_FILE
411 .include CHECK_RCPT_LOCAL_ACL_FILE
412 .endif
413
414 accept
415     domains = +relay_to_domains
416     endpass
417     verify = recipient
418
419 accept
420
421 acl_check_data:
422
423 .ifdef CHECK_DATA_VERIFY_HEADER_SYNTAX
424 deny
425     message = Message headers fail syntax check
426     !acl = acl_local_deny_exceptions
427     !verify = header_syntax
428 .endif
429
430 .ifdef CHECK_DATA_VERIFY_HEADER_SENDER
431 deny
432     message = No verifiable sender address in message headers
433     !acl = acl_local_deny_exceptions
434     !verify = header_sender
435 .endif
436
437 .ifdef CHECK_DATA_LOCAL_ACL_FILE
438 .include CHECK_DATA_LOCAL_ACL_FILE
439 .endif
440
441 accept
442
443 begin routers
```

APPENDIX E. EMAIL CONFIGURATION

```
442
443 .ifdef MAIN_ALLOW_DOMAIN_LITERALS
444 domain_literal:
445     debug_print = "R: domain_literal for $local_part@$domain"
446     driver = ipliteral
447     domains = ! +local_domains
448     transport = remote_smtp
449 .endif
450
451 hubbed_hosts:
452     debug_print = "R: hubbed_hosts for $domain"
453     driver = manualroute
454     domains = "${if exists{CONFFDIR/hubbed_hosts}\
455                 {partial-lsearch;CONFFDIR/hubbed_hosts}\
456                 fail}"
457     same_domain_copy_routing = yes
458     route_data = ${lookup{$domain}partial-lsearch{CONFFDIR/hubbed_hosts}}
459     transport = remote_smtp
460
461 .ifdef DCconfig_internet
462
463 dnslookup_relay_to_domains:
464     debug_print = "R: dnslookup_relay_to_domains for $local_part@$domain"
465     driver = dnslookup
466     domains = ! +local_domains : +relay_to_domains
467     transport = remote_smtp
468     same_domain_copy_routing = yes
469     no_more
470
471 dnslookup:
472     debug_print = "R: dnslookup for $local_part@$domain"
473     driver = dnslookup
474     domains = ! +local_domains
475     transport = remote_smtp
476     same_domain_copy_routing = yes
477     ignore_target_hosts = 0.0.0.0 : 127.0.0.0/8 : 192.168.0.0/16 :\
478                           172.16.0.0/12 : 10.0.0.0/8 : 169.254.0.0/16 :\
479                           255.255.255.255
480     no_more
481
482 .endif
483
484 .ifdef DCconfig_local
485 nonlocal:
486     debug_print = "R: nonlocal for $local_part@$domain"
487     driver = redirect
488     domains = ! +local_domains
```

APPENDIX E. EMAIL CONFIGURATION

```
489     allow_fail
490     data = :fail: Mailing to remote domains not supported
491     no_more
492
493 .endif
494
495 .ifdef DCconfig_smarthost DCconfig_satellite
496
497 smarthost:
498     debug_print = "R: smarthost for $local_part@$domain"
499     driver = manualroute
500     domains = ! +local_domains
501     transport = remote_smtp_smarthost
502     route_list = * DCsmarthost byname
503     host_find_failed = ignore
504     same_domain_copy_routing = yes
505     no_more
506
507 .endif
508
509 COND_LOCAL_SUBMITTER = "\
510     ${if match_ip{$sender_host_address}{:@[]}\
511         {1}{0}\
512     }"
513
514 real_local:
515     debug_print = "R: real_local for $local_part@$domain"
516     driver = accept
517     domains = +local_domains
518     condition = COND_LOCAL_SUBMITTER
519     local_part_prefix = real-
520     check_local_user
521     transport = LOCAL_DELIVERY
522
523 system_aliases:
524     debug_print = "R: system_aliases for $local_part@$domain"
525     driver = redirect
526     domains = +local_domains
527     allow_fail
528     allow_defer
529     data = ${lookup{$local_part}lsearch{/etc/aliases}}
530     .ifdef SYSTEM_ALIASES_USER
531     user = SYSTEM_ALIASES_USER
532     .endif
533     .ifdef SYSTEM_ALIASES_GROUP
534     group = SYSTEM_ALIASES_GROUP
535     .endif
```

APPENDIX E. EMAIL CONFIGURATION

```
536 .ifdef SYSTEM_ALIASES_FILE_TRANSPORT
537 file_transport = SYSTEM_ALIASES_FILE_TRANSPORT
538 .endif
539 .ifdef SYSTEM_ALIASES_PIPE_TRANSPORT
540 pipe_transport = SYSTEM_ALIASES_PIPE_TRANSPORT
541 .endif
542 .ifdef SYSTEM_ALIASES_DIRECTORY_TRANSPORT
543 directory_transport = SYSTEM_ALIASES_DIRECTORY_TRANSPORT
544 .endif
545
546 .ifdef DCconfig_satellite
547 hub_user:
548     debug_print = "R: hub_user for $local_part@$domain"
549     driver = redirect
550     domains = +local_domains
551     data = ${local_part}@DCreadhost
552     check_local_user
553
554 hub_user_smarthost:
555     debug_print = "R: hub_user_smarthost for $local_part@$domain"
556     driver = manualroute
557     domains = DCreadhost
558     transport = remote_smtp_smarthost
559     route_list = * DCsmarthost byname
560     host_find_failed = ignore
561     same_domain_copy_routing = yes
562     check_local_user
563 .endif
564
565 userforward:
566     debug_print = "R: userforward for $local_part@$domain"
567     driver = redirect
568     domains = +local_domains
569     check_local_user
570     file = $home/.forward
571     require_files = $local_part:$home/.forward
572     no_verify
573     no_expn
574     check_ancestor
575     allow_filter
576     forbid_smtp_code = true
577     directory_transport = address_directory
578     file_transport = address_file
579     pipe_transport = address_pipe
580     reply_transport = address_reply
581     skip_syntax_errors
582     syntax_errors_to = real-$local_part@$domain
```

APPENDIX E. EMAIL CONFIGURATION

```
583 syntax_errors_text = \  
584     This is an automatically generated message. An error has\  
585     been found in your .forward file. Details of the error are\  
586     reported below. While this error persists, you will receive\  
587     a copy of this message for every message that is addressed\  
588     to you. If your .forward file is a filter file, or if it is\  
589     a non-filter file containing no valid forwarding addresses,\  
590     a copy of each incoming message will be put in your normal\  
591     mailbox. If a non-filter file contains at least one valid\  
592     forwarding address, forwarding to the valid addresses will\  
593     happen, and those will be the only deliveries that occur.  
594  
595 procmail:  
596     debug_print = "R: procmail for $local_part@$domain"  
597     driver = accept  
598     domains = +local_domains  
599     check_local_user  
600     transport = procmail_pipe  
601     require_files = ${local_part}:\  
602                     ${if exists{/etc/procmailrc}}\  
603                     {/etc/procmailrc}${home}/.procmailrc}:\  
604                     +usr/bin/procmail  
605     no_verify  
606     no_expn  
607  
608 maildrop:  
609     debug_print = "R: maildrop for $local_part@$domain"  
610     driver = accept  
611     domains = +local_domains  
612     check_local_user  
613     transport = maildrop_pipe  
614     require_files = ${local_part}:${home}/.mailfilter:+usr/bin/maildrop  
615     no_verify  
616     no_expn  
617  
618 .ifndef FIRST_USER_ACCOUNT_UID  
619 FIRST_USER_ACCOUNT_UID = 0  
620 .endif  
621  
622 .ifndef DEFAULT_SYSTEM_ACCOUNT_ALIAS  
623 DEFAULT_SYSTEM_ACCOUNT_ALIAS = :fail: no mail to system accounts  
624 .endif  
625  
626 COND_SYSTEM_USER_AND_REMOTE_SUBMITTER = "\  
627     ${if and{${! match_ip{$sender_host_address}{:@[]}}\  
628     {<{$local_user_uid}{FIRST_USER_ACCOUNT_UID}}}  
629     {1}{0}}\
```


APPENDIX E. EMAIL CONFIGURATION

```
630     }"
631
632 lowuid_aliases:
633     debug_print = "R: lowuid_aliases for $local_part@$domain (UID
634         $local_user_uid)"
635     check_local_user
636     driver = redirect
637     allow_fail
638     domains = +local_domains
639     condition = COND_SYSTEM_USER_AND_REMOTE_SUBMITTER
640     data = ${if exists{CONFFDIR/lowuid-aliases}\
641         ${lookup{$local_part}lsearch{CONFFDIR/lowuid-aliases}\
642         ${value}{DEFAULT_SYSTEM_ACCOUNT_ALIAS}}}\
643         {DEFAULT_SYSTEM_ACCOUNT_ALIAS}}
644
645 local_user:
646     debug_print = "R: local_user for $local_part@$domain"
647     driver = accept
648     domains = +local_domains
649     check_local_user
650     local_parts = ! root
651     transport = LOCAL_DELIVERY
652     cannot_route_message = Unknown user
653
654 mail4root:
655     debug_print = "R: mail4root for $local_part@$domain"
656     driver = redirect
657     domains = +local_domains
658     data = /var/mail/mail
659     file_transport = address_file
660     local_parts = root
661     user = mail
662     group = mail
663
664 begin transports
665
666 .ifdef HIDE_MAILNAME
667 REMOTE_SMTP_HEADERS_REWRITE=*@+local_domains $1@DCreadhost frs : *
668     @ETC_MAILNAME $1@DCreadhost frs
669 REMOTE_SMTP_RETURN_PATH=${if match_domain{$sender_address_domain}{+
670     local_domains}${sender_address_local_part}@DCreadhost}${if match_domain{
671     $sender_address_domain}{ETC_MAILNAME}${sender_address_local_part}
672     @DCreadhost}fail}}
673 .endif
674
675 .ifdef REMOTE_SMTP_HELO_FROM_DNS
676 .ifdef REMOTE_SMTP_HELO_DATA
```

APPENDIX E. EMAIL CONFIGURATION

```
672 REMOTE_SMTP_HELO_DATA==${lookup dnsdb {ptr=$sending_ip_address}{$value}{  
    $primary_hostname}}  
673 .else  
674 REMOTE_SMTP_HELO_DATA=${lookup dnsdb {ptr=$sending_ip_address}{$value}{  
    $primary_hostname}}  
675 .endif  
676 .endif  
677  
678 address_file:  
679     debug_print = "T: address_file for $local_part@$domain"  
680     driver = appendfile  
681     delivery_date_add  
682     envelope_to_add  
683     return_path_add  
684  
685 address_pipe:  
686     debug_print = "T: address_pipe for $local_part@$domain"  
687     driver = pipe  
688     return_fail_output  
689  
690 address_reply:  
691     debug_print = "T: autoreply for $local_part@$domain"  
692     driver = autoreply  
693  
694 mail_spool:  
695     debug_print = "T: appendfile for $local_part@$domain"  
696     driver = appendfile  
697     file = /var/mail/$local_part  
698     delivery_date_add  
699     envelope_to_add  
700     return_path_add  
701     group = mail  
702     mode = 0660  
703     mode_fail_narrower = false  
704  
705 maildir_home:  
706     debug_print = "T: maildir_home for $local_part@$domain"  
707     driver = appendfile  
708     .ifdef MAILDIR_HOME_MAILDIR_LOCATION  
709     directory = MAILDIR_HOME_MAILDIR_LOCATION  
710     .else  
711     directory = $home/Maildir  
712     .endif  
713     .ifdef MAILDIR_HOME_CREATE_DIRECTORY  
714     create_directory  
715     .endif  
716     .ifdef MAILDIR_HOME_CREATE_FILE
```

APPENDIX E. EMAIL CONFIGURATION

```
717 create_file = MAILDIR_HOME_CREATE_FILE
718 .endif
719 delivery_date_add
720 envelope_to_add
721 return_path_add
722 maildir_format
723 .ifdef MAILDIR_HOME_DIRECTORY_MODE
724 directory_mode = MAILDIR_HOME_DIRECTORY_MODE
725 .else
726 directory_mode = 0700
727 .endif
728 .ifdef MAILDIR_HOME_MODE
729 mode = MAILDIR_HOME_MODE
730 .else
731 mode = 0600
732 .endif
733 mode_fail_narrower = false
734
735 maildrop_pipe:
736 debug_print = "T: maildrop_pipe for $local_part@$domain"
737 driver = pipe
738 path = "/bin:/usr/bin:/usr/local/bin"
739 command = "/usr/bin/maildrop"
740 return_path_add
741 delivery_date_add
742 envelope_to_add
743
744 procmail_pipe:
745 debug_print = "T: procmail_pipe for $local_part@$domain"
746 driver = pipe
747 path = "/bin:/usr/bin:/usr/local/bin"
748 command = "/usr/bin/procmail"
749 return_path_add
750 delivery_date_add
751 envelope_to_add
752
753 remote_smtp:
754 debug_print = "T: remote_smtp for $local_part@$domain"
755 driver = smtp
756 .ifdef REMOTE_SMTP_HOSTS_AVOID_TLS
757 hosts_avoid_tls = REMOTE_SMTP_HOSTS_AVOID_TLS
758 .endif
759 .ifdef REMOTE_SMTP_HEADERS_REWRITE
760 headers_rewrite = REMOTE_SMTP_HEADERS_REWRITE
761 .endif
762 .ifdef REMOTE_SMTP_RETURN_PATH
763 return_path = REMOTE_SMTP_RETURN_PATH
```

APPENDIX E. EMAIL CONFIGURATION

```
764 .endif
765 .ifdef REMOTE_SMTP_HELO_DATA
766     helo_data=REMOTE_SMTP_HELO_DATA
767 .endif
768 .ifdef DKIM_DOMAIN
769 dkim_domain = DKIM_DOMAIN
770 .endif
771 .ifdef DKIM_SELECTOR
772 dkim_selector = DKIM_SELECTOR
773 .endif
774 .ifdef DKIM_PRIVATE_KEY
775 dkim_private_key = DKIM_PRIVATE_KEY
776 .endif
777 .ifdef DKIM_CANON
778 dkim_canon = DKIM_CANON
779 .endif
780 .ifdef DKIM_STRICT
781 dkim_strict = DKIM_STRICT
782 .endif
783 .ifdef DKIM_SIGN_HEADERS
784 dkim_sign_headers = DKIM_SIGN_HEADERS
785 .endif
786 .ifdef TLS_DH_MIN_BITS
787 tls_dh_min_bits = TLS_DH_MIN_BITS
788 .endif
789 .ifdef REMOTE_SMTP_TLS_CERTIFICATE
790 tls_certificate = REMOTE_SMTP_TLS_CERTIFICATE
791 .endif
792 .ifdef REMOTE_SMTP_PRIVATEKEY
793 tls_privatekey = REMOTE_SMTP_PRIVATEKEY
794 .endif
795
796 remote_smtp_smarthost:
797     debug_print = "T: remote_smtp_smarthost for $local_part@$domain"
798     driver = smtp
799     hosts_try_auth = <; ${if exists{CONFDIR/passwd.client} \
800         {\
801             ${lookup{$host}nwildlsearch{CONFDIR/passwd.client}{$host_address}}\
802         }\
803         {} \
804     }
805 .ifdef REMOTE_SMTP_SMARTHOST_HOSTS_AVOID_TLS
806     hosts_avoid_tls = REMOTE_SMTP_SMARTHOST_HOSTS_AVOID_TLS
807 .endif
808 .ifdef REMOTE_SMTP_HEADERS_REWRITE
809     headers_rewrite = REMOTE_SMTP_HEADERS_REWRITE
810 .endif
```

APPENDIX E. EMAIL CONFIGURATION

```
811 .ifdef REMOTE_SMTP_RETURN_PATH
812     return_path = REMOTE_SMTP_RETURN_PATH
813 .endif
814 .ifdef REMOTE_SMTP_HELO_DATA
815     helo_data=REMOTE_SMTP_HELO_DATA
816 .endif
817 .ifdef TLS_DH_MIN_BITS
818     tls_dh_min_bits = TLS_DH_MIN_BITS
819 .endif
820 .ifdef REMOTE_SMTP_SMARTHOST_TLS_CERTIFICATE
821     tls_certificate = REMOTE_SMTP_SMARTHOST_TLS_CERTIFICATE
822 .endif
823 .ifdef REMOTE_SMTP_SMARTHOST_PRIVATEKEY
824     tls_privatekey = REMOTE_SMTP_SMARTHOST_PRIVATEKEY
825 .endif
826
827 address_directory:
828     debug_print = "T: address_directory for $local_part@$domain"
829     driver = appendfile
830     delivery_date_add
831     envelope_to_add
832     return_path_add
833     check_string = ""
834     escape_string = ""
835     maildir_format
836
837 begin retry
838
839 *                *                F,2h,15m; G,16h,1h,1.5; F,4d,6h
840
841 begin rewrite
842
843 .ifndef NO_EAA_REWRITE_REWRITE
844 *@+local_domains "${lookup{${local_part}}lsearch{/etc/email-addresses}\
845                 {$value}fail}" Ffrs
846 *@ETC_MAILNAME "${lookup{${local_part}}lsearch{/etc/email-addresses}\
847                 {$value}fail}" Ffrs
848 .endif
849
850 begin authenticators
851
852 cram_md5:
853     driver = cram_md5
854     public_name = CRAM-MD5
855     client_name = ${extract{1}{:}}${lookup{$host}nwildlsearch{CONFDIR/passwd.
856                 client}{$value}fail}}
856     client_secret = ${extract{2}{:}}${lookup{$host}nwildlsearch{CONFDIR/passwd.
```

APPENDIX E. EMAIL CONFIGURATION

```
client}{${value}fail}}}  
857  
858 PASSWDLINE=${sg\  
859     ${lookup{$host}nwildlsearch{CONFDIR/passwd.client}{${value}fail  
    }\  
860     }\  
861     {\N[\^]\N\  
862     {\^}\  
863     }  
864  
865 plain:  
866     driver = plaintext  
867     public_name = PLAIN  
868 .ifndef AUTH_CLIENT_ALLOW_NOTLS_PASSWORDS  
869     client_send = "<; ${if !eq{$tls_out_cipher}{}\  
870         ${extract{1}{:}{PASSWDLINE}}\  
871         ${sg{PASSWDLINE}{\N([~:]+:)(.*)\N}{\\$2}}\  
872     }fail}"  
873 .else  
874     client_send = "<; ${extract{1}{:}{PASSWDLINE}}\  
875     ${sg{PASSWDLINE}{\N([~:]+:)(.*)\N}{\\$2}}"  
876 .endif  
877  
878 login:  
879     driver = plaintext  
880     public_name = LOGIN  
881 .ifndef AUTH_CLIENT_ALLOW_NOTLS_PASSWORDS  
882     client_send = "<; ${if and\  
883         {!eq{$tls_out_cipher}{}}\  
884         {!eq{PASSWDLINE}{}}\  
885     }\  
886     {}fail\  
887     ; ${extract{1}{:}{PASSWDLINE}}\  
888     ; ${sg{PASSWDLINE}{\N([~:]+:)(.*)\N}{\\$2}}"  
889 .else  
890     client_send = "<; ${if !eq{PASSWDLINE}{}\  
891         {}fail\  
892     ; ${extract{1}{:}{PASSWDLINE}}\  
893     ; ${sg{PASSWDLINE}{\N([~:]+:)(.*)\N}{\\$2}}"  
894 .endif
```